

Lab Observation Book


R V Abhishek

2025-08-13

B.M.S. COLLEGE OF ENGINEERING

(Autonomous College under VTU)

Bull Temple Road, Basavangudi, Bangalore - 560019



Lab Observation

ON

Programming with R

Submitted by

R V Abhishek(1BM23CD047)

in fulfillment of mandatory observation submission for Lab assessment

BACHELOR OF ENGINEERING

in

Computer Science & Engineering (Data Science)

Under the Guidance of

Dr. Kalyan N

Assistant Professor

Department of CSE (Data Science),


B.M.S. College of Engineering

2025-2026

B.M.S. COLLEGE OF ENGINEERING

(Autonomous College under VTU)

Bull Temple Road, Basavangudi, Bangalore - 560019



Laboratory Certificate

This is to certify that Mr./Ms. R V Abhishek has satisfactorily completed the course of experiments in practical **Programming With R** prescribed by the Visvesvaraya Technology University for 5th Semester Bachelor of Engineering course in the laboratory of the college in the year 2024 - 2025

Head of the Department

Staff Incharge of the Batch

| Marks | |
|---------|----------|
| Maximum | Obtained |
| | |

Name of the Candidate: R V Abhishek

Branch: CSE (Data Science)

USN: 1BM23CD047

Date:

Signature of the Candidate

TABLE OF CONTENTS

| SI No | Program No | Date of Execution | Marks | Faculty In-Charge Sign | Page No |
|-------|---|-------------------------|-------|------------------------------|---------|
| 1 | Program to check what type of Triangle given 3 sides, and calculate its area | 2025-08-13 | | | 3–5 |
| 2 | Creating and Manipulating Data Structures | 2025-08-20 | | | 6–14 |
| 3 | Basic Statistical Operations on Open-Source Datasets | 2025-08-26 | | | 14–22 |
| 4 | Data Import, Cleaning, and Export with Titanic Dataset and Adult Income Dataset | 2025-09-09 | | | 23–31 |
| 5 | Advanced Data Manipulation with dplyr and Complex Grouping | 2025-09-16 | | | 32–39 |
| 6 | Data Visualisation with ggplot2 and Customisations | 2025-09-23 | | | 40–45 |
| 7 | Linear and Multiple Regression Analysis with Interaction terms | 2025-10-14 | | | 46 –54 |
| 8 | K-Means Clustering and PCA for Dimensionality Reduction | 2025-10-28 | | | 55–61 |
| 9 | Time Series Analysis using ARIMA and Seasonal Decomposition | 2025-10-28 | | | 62–82 |
| 10 | Interactive Visualization with plotly and Dynamic Reports with RMarkdown | 2025-10-28 | | | 82–89 |

Program - 1

Program to check what type of Triangle given 3 sides, and calculate its area

Date of Execution - 2025-08-13

This R program validates the sides of the triangle (taken as input from the user) and then if valid, calculates the area of the triangle using Heron's formula and checks what type of triangle it is

```
# Validating the triangle
is_valid_triangle <- function(a, b, c) {
  return ((a + b > c) & (b + c > a) & (a + c > b))
}
```

```
# Function to check the type of triangle
triangle_type <- function(a, b, c) {
  if (a == b && b == c) {
    return("Equilateral ")
  } else if (a == b || b == c || a == c) {
    return("Isosceles ")
  } else {
    return("Scalene")
  }
}
```

```
# Calculating Area using Heron's Formula
triangle_area <- function(a, b, c) {
  s <- (a + b + c) / 2 # Semi - perimeter
  # Heron 's formula
  area <- sqrt(s * (s - a) * (s - b) * (s - c))
  return(area)
}
```

```
# Validating inputs
validate_input <- function(x) {
  if (!is.numeric(x) || x <= 0) {
    stop("Error : Input must be a positive number.")
  }
  return(TRUE)
}
```

```
## Main Code Block
```

```
# 1. Defining 3 variables representing the 3 sides of the triangle
```

```
cat("Enter the lengths of the sides of the triangle :\n")
```

Enter the lengths of the sides of the triangle :

```
a <- as.numeric(readline(prompt = "Side A: "))
```

Side A:

```
b <- as.numeric(readline(prompt = "Side B: "))
```

Side B:

```
c <- as.numeric(readline(prompt = "Side C: "))
```

Side C:

```
# 2. Input Validation and implementation of all the functions.
```

```
# Input validation}
```

```
tryCatch ({
```

```
  validate_input(a)
```

```
  validate_input(b)
```

```
  validate_input(c)
```

```
# Check if the inputs form a valid triangle
```

```
if (!is_valid_triangle(a , b , c)) {
```

```
  stop("Error : The given sides do not form a valid triangle.")
```

```
}
```

```
# Determine the type of triangle
```

```
type_of_triangle <- triangle_type(a , b , c)
```

```
cat("The triangle is:", type_of_triangle, "\n")
```

```
# Calculate the area of the triangle
```

```
area_of_triangle <- triangle_area(a, b, c)
```

```
cat("The area of the triangle is:", area_of_triangle, "\n")
```

```
}, error = function(e){
```

```
  cat(e$message, "\n")
```

```
})
```

missing value where TRUE/FALSE needed

Sample Output

Enter the lengths of the sides of the triangle:

Side a: 5

Side b: 5

Side c: 8

The triangle is: Isosceles

The area of the triangle is: 12

Enter the lengths of the sides of the triangle:**

Side a: 1

Side b: 2

Side c: 8

Error: The given sides do not form a valid triangle.

Program - 2

Creating and Manipulating Data Structures

Date of Execution - 2025-08-20

Objective - This program evaluates the student's understanding of different data structures (vectors, matrices, lists, and data frames) in R and how to manipulate them.

1. Create a vector of random numbers and apply operations such as sorting and searching

```
set.seed(42) # For reproducibility
random_vector <- runif(20, min = 1, max = 100)
cat("Original random vector:\n")
```

Original random vector:

```
print(random_vector)
```

```
[1] 91.56580 93.77047 29.32781 83.21431 64.53281 52.39050 73.92224
[8] 14.33199 66.04224 70.80141 46.31644 72.19211 93.53255 26.28745
[15] 46.76699 94.06144 97.84442 12.63125 48.02471 56.47294
```

```
# Sort the vector
sorted_vector <- sort(random_vector)
cat("Sorted vector:\n")
```

Sorted vector:

```
print(sorted_vector)
```

```
[1] 12.63125 14.33199 26.28745 29.32781 46.31644 46.76699 48.02471
[8] 52.39050 56.47294 64.53281 66.04224 70.80141 72.19211 73.92224
[15] 83.21431 91.56580 93.53255 93.77047 94.06144 97.84442
```

```
# Search for a specific value (check if a number is present)
search_value <- 50
is_value_present <- any(random_vector == search_value)
cat("Is", search_value, "present in the vector?", is_value_present, "\n")
```

Is 50 present in the vector? FALSE


```
# Find values in the vector greater than 60
```

```
values_greater_than_60 <- random_vector[random_vector > 60]  
cat("Values greater than 60:\n")
```

Values greater than 60:

```
print(values_greater_than_60)
```

```
[1] 91.56580 93.77047 83.21431 64.53281 73.92224 66.04224 70.80141  
[8] 72.19211 93.53255 94.06144 97.84442
```

```
# 2. Convert the vector into a matrix and perform matrix multiplication
```

```
# Create a 4x5 matrix from the vector
```

```
matrix_from_vector <- matrix(random_vector, nrow = 4, ncol = 5)  
cat("Matrix from vector:\n")
```

Matrix from vector:

```
print(matrix_from_vector)
```

```
      [,1]      [,2]      [,3]      [,4]      [,5]  
[1,] 91.56580 64.53281 66.04224 93.53255 97.84442  
[2,] 93.77047 52.39050 70.80141 26.28745 12.63125  
[3,] 29.32781 73.92224 46.31644 46.76699 48.02471  
[4,] 83.21431 14.33199 72.19211 94.06144 56.47294
```

```
# Perform matrix multiplication (matrix with its transpose)
```

```
matrix_transpose <- t(matrix_from_vector)  
matrix_multiplication_result <- matrix_from_vector %*% matrix_transpose  
cat("Matrix multiplication result:\n")
```

Matrix multiplication result:

```
print(matrix_multiplication_result)
```

```
      [,1]      [,2]      [,3]      [,4]  
[1,] 35232.22 20337.59 19587.86 27635.57  
[2,] 20337.59 17401.08 11738.17 16851.17  
[3,] 19587.86 11738.17 12963.36 13954.70  
[4,] 27635.57 16851.17 13954.70 24378.48
```

```
# Element-wise matrix multiplication (Hadamard product)
elementwise_multiplication_result <- matrix_from_vector * matrix_from_vector
cat("Element-wise matrix multiplication result:\n")
```

Element-wise matrix multiplication result:

```
print(elementwise_multiplication_result)
```

```
      [,1]      [,2]      [,3]      [,4]      [,5]
[1,] 8384.2954 4164.483 4361.577 8748.3384 9573.5298
[2,] 8792.9003 2744.764 5012.840  691.0302  159.5484
[3,]  860.1207 5464.498 2145.212 2187.1513 2306.3729
[4,] 6924.6222  205.406 5211.701 8847.5541 3189.1932
```

```
# 3. Create a list containing different types of elements and perform subsetting
```

```
my_list <- list(
  numbers = random_vector,
  characters = c("A", "B", "C", "D"),
  logical_values = c(TRUE, FALSE, TRUE),
  matrix = matrix_from_vector
)
cat("List:\n")
```

List:

```
print(my_list)
```

\$numbers

```
[1] 91.56580 93.77047 29.32781 83.21431 64.53281 52.39050 73.92224
[8] 14.33199 66.04224 70.80141 46.31644 72.19211 93.53255 26.28745
[15] 46.76699 94.06144 97.84442 12.63125 48.02471 56.47294
```

\$characters

```
[1] "A" "B" "C" "D"
```

\$logical_values

```
[1] TRUE FALSE TRUE
```

\$matrix

```
      [,1]      [,2]      [,3]      [,4]      [,5]
```

```
[1,] 91.56580 64.53281 66.04224 93.53255 97.84442
[2,] 93.77047 52.39050 70.80141 26.28745 12.63125
[3,] 29.32781 73.92224 46.31644 46.76699 48.02471
[4,] 83.21431 14.33199 72.19211 94.06144 56.47294
```

```
# Subsetting the list (extracting numeric and logical parts)
```

```
subset_numeric <- my_list$numbers
cat("Subset (numeric part of the list):\n")
```

```
Subset (numeric part of the list):
```

```
str(subset_numeric)
```

```
num [1:20] 91.6 93.8 29.3 83.2 64.5 ...
```

```
cat("\n")
```

```
subset_logical <- my_list$logical_values
cat("Subset (logical part of the list):\n", subset_logical, "\n")
```

```
Subset (logical part of the list):
```

```
TRUE FALSE TRUE
```

```
# Modify elements in the list (replace the second character with "Z")
```

```
my_list$characters[2] <- "Z"
cat("Modified list of characters:\n", my_list$characters, "\n")
```

```
Modified list of characters:
```

```
A Z C D
```

```
# Apply a function to the numeric part of the list
```

```
# (e.g., calculate the square of the numbers)
```

```
squared_numbers <- my_list$numbers ^ 2
cat("Squared numbers:\n")
```

```
Squared numbers:
```

```
str(squared_numbers)
```

```
num [1:20] 8384 8793 860 6925 4164 ...
```

```
cat("\n")
```

```
# 4. Create a data frame and perform operations such as filtering,  
# summarizing, and handling missing values
```

```
# Create a data frame
```

```
df <- data.frame(  
  ID = 1:20,  
  Age = sample(18:65, 20, replace = TRUE),  
  Score = runif(20, min = 50, max = 100),  
  Passed = sample(c(TRUE, FALSE), 20, replace = TRUE)  
)  
cat("Data frame:\n")
```

Data frame:

```
print(df)
```

| | ID | Age | Score | Passed |
|----|----|-----|----------|--------|
| 1 | 1 | 64 | 71.78858 | FALSE |
| 2 | 2 | 20 | 51.87155 | FALSE |
| 3 | 3 | 58 | 98.67700 | FALSE |
| 4 | 4 | 42 | 71.58756 | FALSE |
| 5 | 5 | 44 | 97.87883 | FALSE |
| 6 | 6 | 53 | 94.38775 | FALSE |
| 7 | 7 | 54 | 81.99894 | TRUE |
| 8 | 8 | 48 | 98.54833 | FALSE |
| 9 | 9 | 62 | 80.94191 | TRUE |
| 10 | 10 | 22 | 66.67136 | FALSE |
| 11 | 11 | 37 | 67.33741 | TRUE |
| 12 | 12 | 51 | 69.92427 | FALSE |
| 13 | 13 | 45 | 89.23464 | FALSE |
| 14 | 14 | 57 | 51.94682 | FALSE |
| 15 | 15 | 20 | 87.43977 | FALSE |
| 16 | 16 | 50 | 83.86384 | TRUE |
| 17 | 17 | 59 | 58.56322 | TRUE |
| 18 | 18 | 41 | 63.05440 | TRUE |
| 19 | 19 | 47 | 75.72065 | TRUE |
| 20 | 20 | 60 | 83.78036 | FALSE |

```
# Filter the data frame (rows where Age > 30 and Score > 70)
filtered_df <- subset(df, Age > 30 & Score > 70)
cat("Filtered data frame (Age > 30 and Score > 70):\n")
```

Filtered data frame (Age > 30 and Score > 70):

```
print(filtered_df)
```

| | ID | Age | Score | Passed |
|----|----|-----|----------|--------|
| 1 | 1 | 64 | 71.78858 | FALSE |
| 3 | 3 | 58 | 98.67700 | FALSE |
| 4 | 4 | 42 | 71.58756 | FALSE |
| 5 | 5 | 44 | 97.87883 | FALSE |
| 6 | 6 | 53 | 94.38775 | FALSE |
| 7 | 7 | 54 | 81.99894 | TRUE |
| 8 | 8 | 48 | 98.54833 | FALSE |
| 9 | 9 | 62 | 80.94191 | TRUE |
| 13 | 13 | 45 | 89.23464 | FALSE |
| 16 | 16 | 50 | 83.86384 | TRUE |
| 19 | 19 | 47 | 75.72065 | TRUE |
| 20 | 20 | 60 | 83.78036 | FALSE |

```
# Calculate mean, sum, and variance of numerical columns (Age and Score)
mean_age <- mean(df$Age)
sum_age <- sum(df$Age)
var_age <- var(df$Age)

mean_score <- mean(df$Score)
sum_score <- sum(df$Score)
var_score <- var(df$Score)

cat("Summary statistics for Age column:\n")
```

Summary statistics for Age column:

```
cat("Mean Age:", mean_age, "\n")
```

Mean Age: 46.7

```
cat("Sum of Age:", sum_age, "\n")
```

Sum of Age: 934

```
cat("Variance of Age:", var_age, "\n")
```

Variance of Age: 179.6947

```
cat("Summary statistics for Score column:\n")
```

Summary statistics for Score column:

```
cat("Mean Score:", mean_score, "\n")
```

Mean Score: 77.26086

```
cat("Sum of Score:", sum_score, "\n")
```

Sum of Score: 1545.217

```
cat("Variance of Score:", var_score, "\n")
```

Variance of Score: 219.2162

```
# 5. Handling missing values in the data frame
```

```
# Introduce some NA values in the Score column
```

```
df$Score[sample(1:20, 5)] <- NA
```

```
cat("Data frame with missing values:\n")
```

Data frame with missing values:

```
print(df)
```

| | ID | Age | Score | Passed |
|---|----|-----|----------|--------|
| 1 | 1 | 64 | 71.78858 | FALSE |
| 2 | 2 | 20 | 51.87155 | FALSE |
| 3 | 3 | 58 | 98.67700 | FALSE |
| 4 | 4 | 42 | NA | FALSE |
| 5 | 5 | 44 | 97.87883 | FALSE |
| 6 | 6 | 53 | 94.38775 | FALSE |

```

7  7  54 81.99894  TRUE
8  8  48 98.54833  FALSE
9  9  62      NA   TRUE
10 10 22      NA  FALSE
11 11 37 67.33741  TRUE
12 12 51      NA  FALSE
13 13 45      NA  FALSE
14 14 57 51.94682  FALSE
15 15 20 87.43977  FALSE
16 16 50 83.86384  TRUE
17 17 59 58.56322  TRUE
18 18 41 63.05440  TRUE
19 19 47 75.72065  TRUE
20 20 60 83.78036  FALSE

```

```

# Replace NA values with the mean of the Score column
df$Score[is.na(df$Score)] <- mean(df$Score, na.rm = TRUE)
cat("Data frame after imputation of missing values:\n")

```

Data frame after imputation of missing values:

```
print(df)
```

```

      ID Age   Score Passed
1     1  64 71.78858  FALSE
2     2  20 51.87155  FALSE
3     3  58 98.67700  FALSE
4     4  42 77.79050  FALSE
5     5  44 97.87883  FALSE
6     6  53 94.38775  FALSE
7     7  54 81.99894   TRUE
8     8  48 98.54833  FALSE
9     9  62 77.79050   TRUE
10    10 22 77.79050  FALSE
11    11 37 67.33741   TRUE
12    12 51 77.79050  FALSE
13    13 45 77.79050  FALSE
14    14 57 51.94682  FALSE
15    15 20 87.43977  FALSE
16    16 50 83.86384   TRUE
17    17 59 58.56322   TRUE

```

```
18 18 41 63.05440 TRUE
19 19 47 75.72065 TRUE
20 20 60 83.78036 FALSE
```

```
# Grouping the data by Passed status and calculating group-wise statistics
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

```
filter, lag
```

The following objects are masked from 'package:base':

```
intersect, setdiff, setequal, union
```

```
grouped_stats <- df %>%
  group_by(Passed) %>%
  summarise(
    mean_score = mean(Score, na.rm = TRUE),
    mean_age = mean(Age)
  )
cat("Grouped statistics by Passed status:\n")
```

Grouped statistics by Passed status:

```
print(grouped_stats)
```

```
# A tibble: 2 x 3
  Passed mean_score mean_age
  <lgl>      <dbl>    <dbl>
1 FALSE      80.6      44.9
2 TRUE       72.6      50
```


Program - 3

Basic Statistical Operations on Open-Source Datasets

Date of Execution - 2025-08-26

Objective: This program emphasizes the application of statistical concepts on real-world datasets and visualization of the data.

```
# Load necessary
library(dplyr) # For data manipulation
library(ggplot2) # For visualization
library(moments) # For skewness and kurtosis
library(palmerpenguins) # For Palmer Penguins dataset

data(iris) # Load Iris dataset

data(penguins) # Load Palmer Penguins

# Function to calculate mode
calc_mode <- function(x) {
  return (as.numeric (names (sort (table (x), decreasing = TRUE)) [1] ))
}

# Perform Statistical Analysis on Iris Dataset
print("----- Iris Dataset Analysis -----")

[1] "----- Iris Dataset Analysis -----"

# Mean
iris_mean <- sapply (iris[, 1:4], mean, na.rm = TRUE )
print(paste("Mean of Iris dataset : ", iris_mean))

[1] "Mean of Iris dataset :  5.84333333333333"
[2] "Mean of Iris dataset :  3.05733333333333"
[3] "Mean of Iris dataset :  3.758"
[4] "Mean of Iris dataset :  1.19933333333333"

#Median
iris_median <- sapply(iris[, 1:4], median, na.rm = TRUE )
print(paste("Median of Iris dataset : ", iris_median))
```

```
[1] "Median of Iris dataset : 5.8" "Median of Iris dataset : 3"
[3] "Median of Iris dataset : 4.35" "Median of Iris dataset : 1.3"
```

#Mode

```
iris_mode <- sapply(iris[, 1:4], calc_mode )
print(paste("Mode of Iris dataset : ", iris_mode))
```

```
[1] "Mode of Iris dataset : 5.8" "Mode of Iris dataset : 3"
[3] "Mode of Iris dataset : 4.35" "Mode of Iris dataset : 1.3"
```

#Variance

```
iris_variance <- sapply(iris[, 1:4], var, na.rm = TRUE )
print(paste("Variance of Iris dataset : ", iris_variance))
```

```
[1] "Variance of Iris dataset : 0.685693512304251"
[2] "Variance of Iris dataset : 0.189979418344519"
[3] "Variance of Iris dataset : 3.11627785234899"
[4] "Variance of Iris dataset : 0.581006263982103"
```

#Standard Deviation

```
iris_sd <- sapply(iris[, 1:4], sd, na.rm = TRUE )
print(paste("Standard Deviation of Iris dataset : ", iris_sd))
```

```
[1] "Standard Deviation of Iris dataset : 0.828066127977863"
[2] "Standard Deviation of Iris dataset : 0.435866284936698"
[3] "Standard Deviation of Iris dataset : 1.76529823325947"
[4] "Standard Deviation of Iris dataset : 0.762237668960347"
```

#Skewness

```
iris_skewness <- sapply(iris[, 1:4], skewness, na.rm = TRUE )
print(paste("Skewness of Iris dataset : ", iris_skewness))
```

```
[1] "Skewness of Iris dataset : 0.311753058502296"
[2] "Skewness of Iris dataset : 0.315767106338938"
[3] "Skewness of Iris dataset : -0.272127666456721"
[4] "Skewness of Iris dataset : -0.101934206565599"
```

Hypothesis Testing (t-test) between Sepal.Length of Setosa and Versicolor


```
setosa <- subset(iris, Species == "setosa")$Sepal.Length
versicolor <- subset(iris, Species == "versicolor")$Sepal.Length
t_test <- t.test(setosa, versicolor)
print(t_test)
```

Welch Two Sample t-test


```
data: setosa and versicolor
t = -10.521, df = 86.538, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -1.1057074 -0.7542926
sample estimates:
mean of x mean of y
 5.006    5.936
```

```
# Visualization of Iris Dataset
# Histogram for Sepal.Length
ggplot(iris, aes(x = Sepal.Length)) +
  geom_histogram(binwidth = 0.3, fill = "blue", color = "black") +
  ggtitle("Histogram of Sepal Length in Iris Dataset")
```

Histogram of Sepal Length in Iris Dataset



```
# Boxplot for Sepal.Length across Species
ggplot(iris, aes(x = Species, y = Sepal.Length, fill = Species)) +
  geom_boxplot() +
  ggtitle("Boxplot of Sepal Length by Species in Iris Dataset")
```



```
print("----- Palmer Penguins Dataset Analysis -----")
```

```
[1] "----- Palmer Penguins Dataset Analysis -----"
```

```
# Remove rows with missing values
```

```
penguins_clean <- na.omit(penguins)
```

```
# Mean
```

```
penguins_mean <- sapply(penguins_clean[, 3:6], mean, na.rm = TRUE)
```

```
print(paste("Mean of Palmer Penguins dataset:", penguins_mean))
```

```
[1] "Mean of Palmer Penguins dataset: 43.9927927927928"
```

```
[2] "Mean of Palmer Penguins dataset: 17.1648648648649"
```

```
[3] "Mean of Palmer Penguins dataset: 200.966966966967"
```

```
[4] "Mean of Palmer Penguins dataset: 4207.05705705706"
```

```
# Median
```

```
penguins_median <- sapply(penguins_clean[, 3:6], median, na.rm = TRUE)
```

```
print(paste("Median of Palmer Penguins dataset:", penguins_median))
```

```
[1] "Median of Palmer Penguins dataset: 44.5"
```

```
[2] "Median of Palmer Penguins dataset: 17.3"
```

```
[3] "Median of Palmer Penguins dataset: 197"  
[4] "Median of Palmer Penguins dataset: 4050"
```

Mode

```
penguins_mode <- sapply(penguins_clean[, 3:6], calc_mode)  
print(paste("Mode of Palmer Penguins dataset:", penguins_mode))
```

```
[1] "Mode of Palmer Penguins dataset: 41.1"  
[2] "Mode of Palmer Penguins dataset: 17"  
[3] "Mode of Palmer Penguins dataset: 190"  
[4] "Mode of Palmer Penguins dataset: 3800"
```

Variance

```
penguins_variance <- sapply(penguins_clean[, 3:6], var, na.rm = TRUE)  
print(paste("Variance of Palmer Penguins dataset:", penguins_variance))
```

```
[1] "Variance of Palmer Penguins dataset: 29.9063334418756"  
[2] "Variance of Palmer Penguins dataset: 3.87788830999674"  
[3] "Variance of Palmer Penguins dataset: 196.441676616375"  
[4] "Variance of Palmer Penguins dataset: 648372.487698542"
```

Standard Deviation

```
penguins_sd <- sapply(penguins_clean[, 3:6], sd, na.rm = TRUE)  
print(paste("Standard Deviation of Palmer Penguins dataset:", penguins_sd))
```

```
[1] "Standard Deviation of Palmer Penguins dataset: 5.46866834264756"  
[2] "Standard Deviation of Palmer Penguins dataset: 1.9692354633199"  
[3] "Standard Deviation of Palmer Penguins dataset: 14.0157652882879"  
[4] "Standard Deviation of Palmer Penguins dataset: 805.215801942897"
```

Skewness

```
penguins_skewness <- sapply(penguins_clean[, 3:6], skewness, na.rm = TRUE)  
print(paste("Skewness of Palmer Penguins dataset:", penguins_skewness))
```

```
[1] "Skewness of Palmer Penguins dataset: 0.0451359779776739"  
[2] "Skewness of Palmer Penguins dataset: -0.149044996398334"  
[3] "Skewness of Palmer Penguins dataset: 0.358523654622741"  
[4] "Skewness of Palmer Penguins dataset: 0.470116171418382"
```

Kurtosis

```
penguins_kurtosis <- sapply(penguins_clean[, 3:6], kurtosis, na.rm = TRUE)  
print(paste("Kurtosis of Palmer Penguins dataset:", penguins_kurtosis))
```

```
[1] "Kurtosis of Palmer Penguins dataset: 2.11182658541194"
[2] "Kurtosis of Palmer Penguins dataset: 2.10341274887238"
[3] "Kurtosis of Palmer Penguins dataset: 2.03516741259049"
[4] "Kurtosis of Palmer Penguins dataset: 2.25951411974012"
```

```
# Hypothesis Testing (t-test) between flipper_length_mm of Adelie and Gentoo species
```

```
adelie <- subset(penguins_clean, species == "Adelie")$flipper_length_mm
gentoo <- subset(penguins_clean, species == "Gentoo")$flipper_length_mm
t_test_penguins <- t.test(adelie, gentoo)
print(t_test_penguins)
```

Welch Two Sample t-test


```
data: adelie and gentoo
t = -33.506, df = 251.35, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -28.72740 -25.53771
sample estimates:
mean of x mean of y
 190.1027  217.2353
```

```
# Visualization of Palmer Penguins Dataset
```

```
# Histogram for flipper_length_mm
```


```
ggplot(penguins_clean, aes(x = flipper_length_mm)) +
  geom_histogram(binwidth = 3, fill = "green", color = "black") +
  ggtitle("Histogram of Flipper Length in Palmer Penguins Dataset")
```

Histogram of Flipper Length in Palmer Penguins Dataset



```
# Boxplot for flipper_length_mm across Species  
ggplot(penguins_clean, aes(x = species, y = flipper_length_mm, fill = species)) +  
  geom_boxplot() +  
  ggtitle("Boxplot of Flipper Length by Species in Palmer Penguins Dataset")
```

Boxplot of Flipper Length by Species in Palmer Penguins Dataset



Program - 4

Data Import, Cleaning, and Export with Titanic Dataset and Adult Income Dataset

Date of Execution - 2025-09-09

Objective: Real World Data Cleaning Processes and emphasis on usage of advanced data wrangling techniques in R.

```
# Load necessary libraries
library(tidyverse)
library(titanic)
library(dplyr)
library(caret)
library(ggcorrplot)

# Load the Titanic dataset
data <- titanic::titanic_train

# Handle the missing data
# Replace missing values in the 'Age' column with the median age
data$Age[is.na(data$Age)] <- median(data$Age, na.rm = TRUE)

# Replace missing values in the 'Embarked' column with the mode
mode_embarked <- as.character(names(sort(table(data$Embarked), decreasing = TRUE)[1]))
data$Embarked[is.na(data$Embarked)] <- mode_embarked

# Define the numeric columns for z-score and correlation calculation
numeric_columns <- c("Age", "SibSp", "Parch", "Fare", "Survived", "Pclass")

# Remove outliers using z-score
z_scores <- as.data.frame(scale(data[, numeric_columns]))

# Identify the rows that have z_scores greater than 3 or less than -3 (outliers)
outlier_rows <- apply(z_scores, 1, function(row) any(abs(row) > 3))

# Filter out Outliers
data_cleaned <- data[!outlier_rows, ]
```

```

# Summarize the dataset before and after cleaning
summary_before <- summary(data)
summary_after <- summary(data_cleaned)

# Calculate Correlation Matrix (fixed)
correlation_matrix <- cor(data_cleaned[, numeric_columns], use = "complete.obs")

# Export cleaned data onto a new CSV file
write.csv(data_cleaned, "titanic_cleaned.csv", row.names = FALSE)

# Display Summaries
print("Summary Before Cleaning:")

```

```
[1] "Summary Before Cleaning:"
```

```
print(summary_before)
```

```

  PassengerId      Survived  Pclass     Name
Min.   : 1.0   Min.   :0.0000   Min.   :1.000   Length:891
1st Qu.:223.5   1st Qu.:0.0000   1st Qu.:2.000   Class  :character
Median :446.0   Median :0.0000   Median :3.000   Mode   :character
Mean   :446.0   Mean    :0.3838   Mean    :2.309
3rd Qu.:668.5   3rd Qu.:1.0000   3rd Qu.:3.000
Max.   :891.0   Max.    :1.0000   Max.    :3.000

   Sex      Age      SibSp      Parch
Length:891   Min.   : 0.42   Min.   :0.000   Min.   :0.0000
Class :character 1st Qu.:22.00   1st Qu.:0.000   1st Qu.:0.0000
Mode  :character Median :28.00   Median :0.000   Median :0.0000
                Mean  :29.36   Mean  :0.523   Mean  :0.3816
                3rd Qu.:35.00   3rd Qu.:1.000   3rd Qu.:0.0000
                Max.   :80.00   Max.   :8.000   Max.   :6.0000

   Ticket      Fare      Cabin
Length:891   Min.   : 0.00   Length:891
Class :character 1st Qu.: 7.91   Class :character
Mode  :character Median :14.45   Mode  :character
                Mean   :32.20
                3rd Qu.:31.00
                Max.   :512.33

   Embarked
Length:891
Class :character

```

```
Mode :character
```

```
print("Summary After Cleaning:")
```

```
[1] "Summary After Cleaning:"
```

```
print(summary_after)
```

| PassengerId | Survived | Pclass | Name |
|---------------|----------------|---------------|------------------|
| Min. : 1.0 | Min. :0.0000 | Min. :1.000 | Length:820 |
| 1st Qu.:226.8 | 1st Qu.:0.0000 | 1st Qu.:2.000 | Class :character |
| Median :446.5 | Median :0.0000 | Median :3.000 | Mode :character |
| Mean :445.7 | Mean :0.3902 | Mean :2.311 | |
| 3rd Qu.:661.2 | 3rd Qu.:1.0000 | 3rd Qu.:3.000 | |
| Max. :891.0 | Max. :1.0000 | Max. :3.000 | |

| Sex | Age | SibSp | Parch |
|------------------|---------------|----------------|----------------|
| Length:820 | Min. : 0.42 | Min. :0.0000 | Min. :0.0000 |
| Class :character | 1st Qu.:23.00 | 1st Qu.:0.0000 | 1st Qu.:0.0000 |
| Mode :character | Median :28.00 | Median :0.0000 | Median :0.0000 |
| | Mean :29.44 | Mean :0.3488 | Mean :0.2549 |
| | 3rd Qu.:35.00 | 3rd Qu.:1.0000 | 3rd Qu.:0.0000 |
| | Max. :66.00 | Max. :3.0000 | Max. :2.0000 |

| Ticket | Fare | Cabin |
|------------------|-----------------|------------------|
| Length:820 | Min. : 0.000 | Length:820 |
| Class :character | 1st Qu.: 7.896 | Class :character |
| Mode :character | Median : 13.000 | Mode :character |
| | Mean : 25.836 | |
| | 3rd Qu.: 27.000 | |
| | Max. :164.867 | |

| Embarked |
|------------------|
| Length:820 |
| Class :character |
| Mode :character |

```
print("Correlation Matrix:")
```

```
[1] "Correlation Matrix:"
```

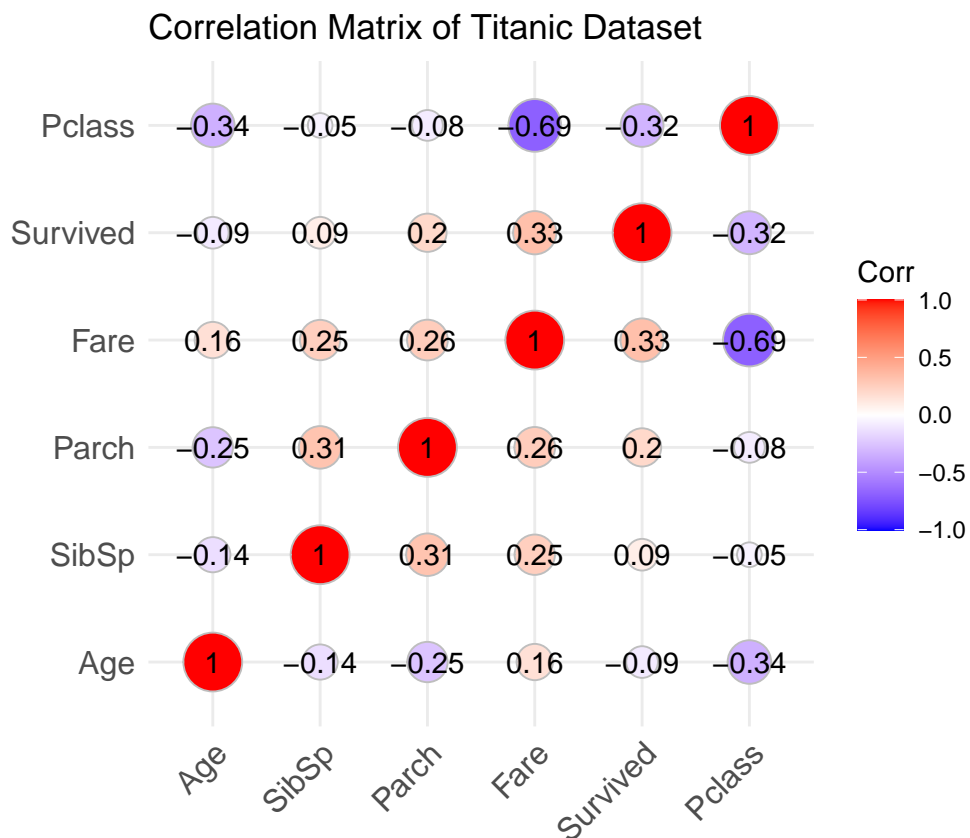
```
print(correlation_matrix)
```

| | Age | SibSp | Parch | Fare | Survived |
|----------|-------------|-------------|------------|------------|-------------|
| Age | 1.00000000 | -0.14391182 | -0.2517719 | 0.1598100 | -0.08602643 |
| SibSp | -0.14391182 | 1.00000000 | 0.3072105 | 0.2472157 | 0.09445934 |
| Parch | -0.25177192 | 0.30721046 | 1.0000000 | 0.2599031 | 0.20107069 |
| Fare | 0.15981001 | 0.24721568 | 0.2599031 | 1.0000000 | 0.33043946 |
| Survived | -0.08602643 | 0.09445934 | 0.2010707 | 0.3304395 | 1.00000000 |
| Pclass | -0.33698055 | -0.05231213 | -0.0783660 | -0.6917198 | -0.32230582 |

| | Pclass |
|----------|-------------|
| Age | -0.33698055 |
| SibSp | -0.05231213 |
| Parch | -0.07836600 |
| Fare | -0.69171982 |
| Survived | -0.32230582 |
| Pclass | 1.00000000 |

```
# Visualize Correlation Matrix (fixed)
```

```
ggcorrplot(correlation_matrix,  
            method = "circle",  
            lab = TRUE) +  
ggtitle("Correlation Matrix of Titanic Dataset")
```



Objective - Data Import, Cleaning, and Export with Adult Income Dataset

```
# Load necessary libraries
library(tidyverse)
library(dplyr)
library(caret)
library(ggcorrplot)

# Load the Adult Income dataset
data <- read.csv("D:/Coding/Coding/Time Series Analysis/Lab 4/adult.data", header = FALSE)

# Assign column names based on the dataset documentation
colnames(data) <- c('age', 'workclass', 'fnlwgt', 'education', 'education_num',
                    'marital_status', 'occupation', 'relationship', 'race', 'sex', 'capital_gain',
                    'capital_loss', 'hours_per_week', 'native_country', 'income')

# Handle missing values represented by '?'
data[data == '?'] <- NA

# Replace categorical missing values with mode
replace_mode <- function(x){
  mode_val <- as.character(names(sort(table(x), decreasing = TRUE)[1]))
}
```

```

x[is.na(x)] <- mode_val
return(x)
}

data <- data %>%
  mutate_if(is.character, replace_mode)

# Replace numeric missing values with median
data <- data %>%
  mutate_if(is.numeric, ~ifelse(is.na(.), median(., na.rm = TRUE), .))

# Define the remove_outliers function
remove_outliers <- function(x){
  z_scores <- scale(x)
  x[abs(z_scores) <= 3]
}

# Remove outliers using z-score
numeric_columns <- sapply(data, is.numeric)

# Apply z-score outlier removal to numeric columns
data_cleaned <- data %>%
  filter(!apply(as.data.frame(scale(data[, numeric_columns])), 1,
    function(row) any(abs(row) > 3)))

# Summarize before and after cleaning
summary_before <- summary(data)
summary_after <- summary(data_cleaned)

# Calculate correlation Matrix
correlation_matrix <- cor(data_cleaned[, numeric_columns], use = "complete.obs")

# Export as CSV
write.csv(data_cleaned, "cleaned_adult_income_data.csv", row.names = FALSE)

# Display Summaries
print("Summary Before Cleaning:")

[1] "Summary Before Cleaning:"

```

```
print(summary_before)
```

| age | workclass | fnlwgt |
|---------------|------------------|-----------------|
| Min. :17.00 | Length:32561 | Min. : 12285 |
| 1st Qu.:28.00 | Class :character | 1st Qu.: 117827 |
| Median :37.00 | Mode :character | Median : 178356 |
| Mean :38.58 | | Mean : 189778 |
| 3rd Qu.:48.00 | | 3rd Qu.: 237051 |
| Max. :90.00 | | Max. :1484705 |

| education | education_num | marital_status |
|------------------|---------------|------------------|
| Length:32561 | Min. : 1.00 | Length:32561 |
| Class :character | 1st Qu.: 9.00 | Class :character |
| Mode :character | Median :10.00 | Mode :character |
| | Mean :10.08 | |
| | 3rd Qu.:12.00 | |
| | Max. :16.00 | |

| occupation | relationship | race |
|------------------|------------------|------------------|
| Length:32561 | Length:32561 | Length:32561 |
| Class :character | Class :character | Class :character |
| Mode :character | Mode :character | Mode :character |

| sex | capital_gain | capital_loss | hours_per_week |
|------------------|--------------|--------------|----------------|
| Length:32561 | Min. : 0 | Min. : 0.0 | Min. : 1.00 |
| Class :character | 1st Qu.: 0 | 1st Qu.: 0.0 | 1st Qu.:40.00 |
| Mode :character | Median : 0 | Median : 0.0 | Median :40.00 |
| | Mean : 1078 | Mean : 87.3 | Mean :40.44 |
| | 3rd Qu.: 0 | 3rd Qu.: 0.0 | 3rd Qu.:45.00 |
| | Max. :99999 | Max. :4356.0 | Max. :99.00 |

| native_country | income |
|------------------|------------------|
| Length:32561 | Length:32561 |
| Class :character | Class :character |
| Mode :character | Mode :character |

```
print("Summary After Cleaning:")
```

```
[1] "Summary After Cleaning:"
```

```
print(summary_after)
```

| age | workclass | fnlwgt |
|------------------|------------------|------------------|
| Min. :17.00 | Length:29828 | Min. : 12285 |
| 1st Qu.:27.00 | Class :character | 1st Qu.:117509 |
| Median :37.00 | Mode :character | Median :177667 |
| Mean :38.14 | | Mean :185193 |
| 3rd Qu.:47.00 | | 3rd Qu.:234279 |
| Max. :79.00 | | Max. :506329 |
| education | education_num | marital_status |
| Length:29828 | Min. : 3.00 | Length:29828 |
| Class :character | 1st Qu.: 9.00 | Class :character |
| Mode :character | Median :10.00 | Mode :character |
| | Mean :10.08 | |
| | 3rd Qu.:12.00 | |
| | Max. :16.00 | |
| occupation | relationship | race |
| Length:29828 | Length:29828 | Length:29828 |
| Class :character | Class :character | Class :character |
| Mode :character | Mode :character | Mode :character |

| sex | capital_gain | capital_loss |
|------------------|------------------|------------------|
| Length:29828 | Min. : 0.0 | Min. : 0.000 |
| Class :character | 1st Qu.: 0.0 | 1st Qu.: 0.000 |
| Mode :character | Median : 0.0 | Median : 0.000 |
| | Mean : 570.2 | Mean : 1.209 |
| | 3rd Qu.: 0.0 | 3rd Qu.: 0.000 |
| | Max. :22040.0 | Max. :1258.000 |
| hours_per_week | native_country | income |
| Min. : 4.0 | Length:29828 | Length:29828 |
| 1st Qu.:40.0 | Class :character | Class :character |
| Median :40.0 | Mode :character | Mode :character |
| Mean :39.9 | | |
| 3rd Qu.:45.0 | | |
| Max. :77.0 | | |

```
print("Correlation Matrix:")
```

```
[1] "Correlation Matrix:"
```



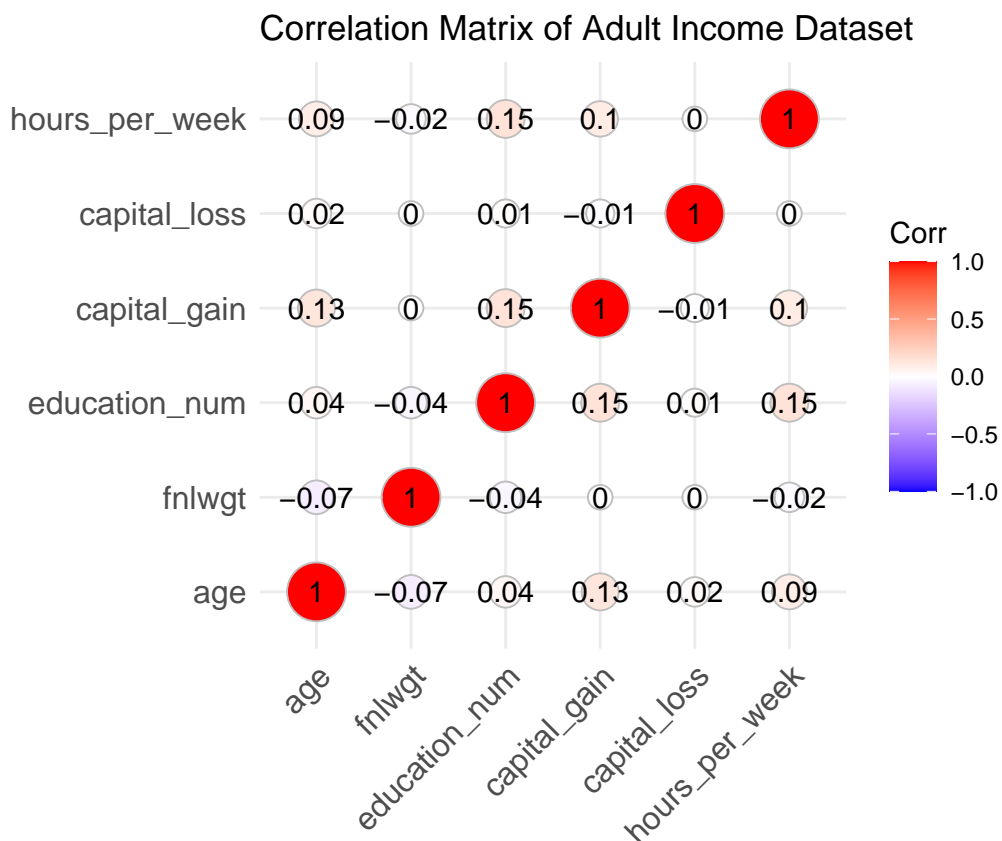
```
print(correlation_matrix)
```

| | age | fnlwgt | education_num | capital_gain |
|----------------|-------------|--------------|---------------|--------------|
| age | 1.00000000 | -0.074427786 | 0.041427102 | 0.131043981 |
| fnlwgt | -0.07442779 | 1.000000000 | -0.037482414 | -0.002378925 |
| education_num | 0.04142710 | -0.037482414 | 1.000000000 | 0.154844283 |
| capital_gain | 0.13104398 | -0.002378925 | 0.154844283 | 1.000000000 |
| capital_loss | 0.02082465 | 0.002583047 | 0.009481359 | -0.009038231 |
| hours_per_week | 0.09219535 | -0.015375555 | 0.150513483 | 0.097209049 |

| | capital_loss | hours_per_week |
|----------------|--------------|----------------|
| age | 0.020824647 | 0.092195352 |
| fnlwgt | 0.002583047 | -0.015375555 |
| education_num | 0.009481359 | 0.150513483 |
| capital_gain | -0.009038231 | 0.097209049 |
| capital_loss | 1.000000000 | -0.003089539 |
| hours_per_week | -0.003089539 | 1.000000000 |

```
# Visualize Correlation Matrix (fixed)
```

```
ggcorrplot(correlation_matrix,
            method = "circle",
            lab = TRUE) +
ggtitle("Correlation Matrix of Adult Income Dataset")
```



Program - 5

Advanced Data Manipulation with dplyr and Complex Grouping

Date of Execution - 2025-09-16

Objective - The goal of this program is to test advanced data manipulation techniques using the dplyr package.

```
## Load necessary libraries
```

```
library(dplyr)
```

```
library(nycflights13)
```

```
library(ggplot2)
```

```
library(zoo)
```

```
# Preview the Star Wars Dataset
```

```
data("starwars")
```

```
head(starwars)
```

```
# A tibble: 6 x 14
```

| | name | height | mass | hair_color | skin_color | eye_color | birth_year | sex |
|---|--------|--------|-------|------------|------------|-----------|------------|-------|
| | <chr> | <int> | <dbl> | <chr> | <chr> | <chr> | <dbl> | <chr> |
| 1 | Luke ~ | 172 | 77 | blond | fair | blue | 19 | male |
| 2 | C-3P0 | 167 | 75 | <NA> | gold | yellow | 112 | none |
| 3 | R2-D2 | 96 | 32 | <NA> | white, bl~ | red | 33 | none |
| 4 | Darth~ | 202 | 136 | none | white | yellow | 41.9 | male |
| 5 | Leia ~ | 150 | 49 | brown | light | brown | 19 | fema~ |
| 6 | Owen ~ | 178 | 120 | brown, gr~ | light | blue | 52 | male |

```
# i 6 more variables: gender <chr>, homeworld <chr>, species <chr>,
```

```
# films <list>, vehicles <list>, starships <list>
```

```
# Select specific columns (name, species, height, mass),
```

```
# filter out rows with missing species or height,
```

```
# and arrange by height in descending order
```

```
starwars_filtered <- starwars %>%
```

```
  select(name, species, height, mass) %>%
```

```
  filter(!is.na(species) & !is.na(height) & height > 100) %>%
```

```
  arrange(desc(height))
```

```
# Display the filtered data
```

```
head(starwars_filtered)
```

```
# A tibble: 6 x 4
```

```

  name      species height  mass
<chr>    <chr>    <int> <dbl>
1 Yarael Poof  Quermian    264    NA
2 Tarfful     Wookiee     234   136
3 Lama Su     Kaminoan    229    88
4 Chewbacca   Wookiee     228   112
5 Roos Tarpals Gungan     224    82
6 Grievous    Kaleesh     216   159

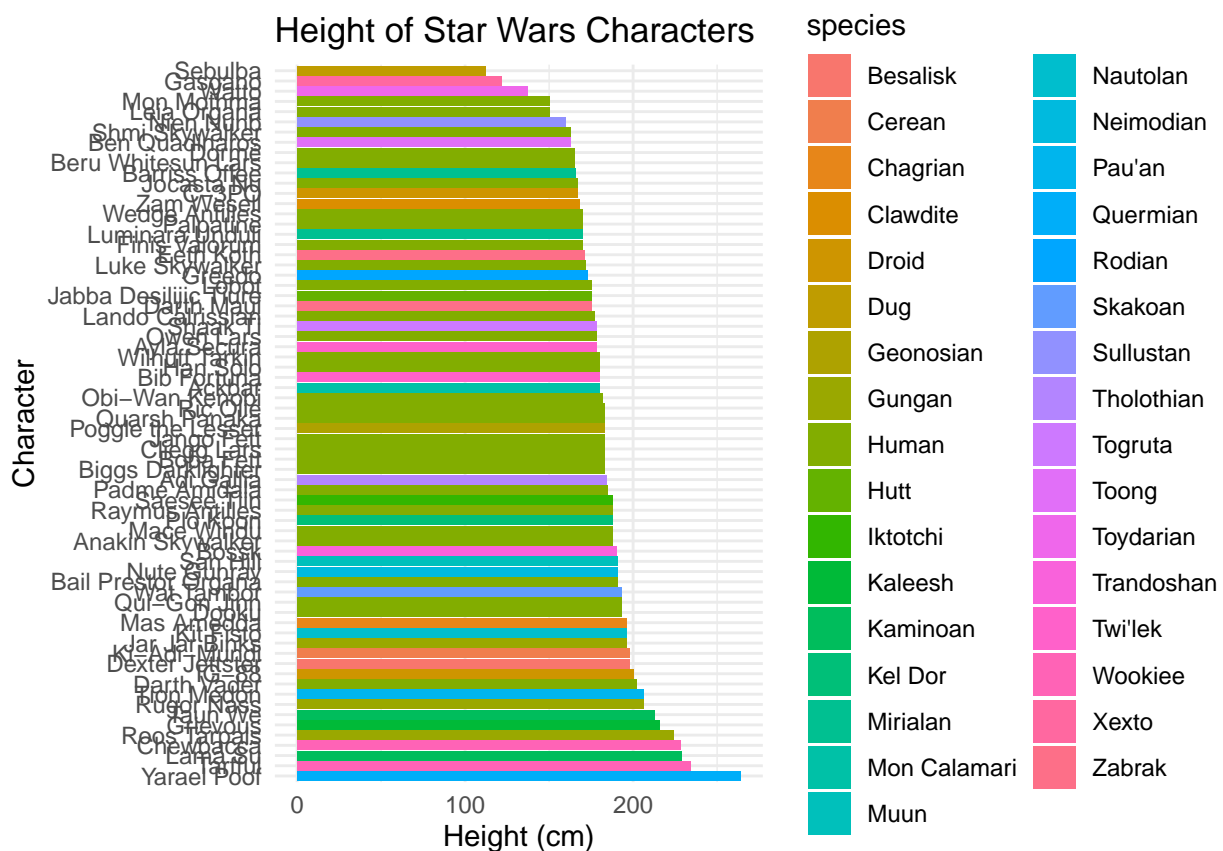
```

```
# Plotting the filtered data
```

```

ggplot(starwars_filtered, aes(x = reorder(name, -height), y = height, fill = species)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  labs(title = "Height of Star Wars Characters",
       x = "Character",
       y = "Height (cm)") +
  theme_minimal()

```



```
# Grouping by species, calculating average height and mass, and counting observation
```

```

species_summary <- starwars %>%
  group_by(species) %>%

```

```

summarise(
  avg_height = mean(height, na.rm = TRUE),
  avg_mass = mean(mass, na.rm = TRUE),
  count = n()
) %>%
  arrange(desc(count))

```

Display the species summary

```
head(species_summary)
```

A tibble: 6 x 4

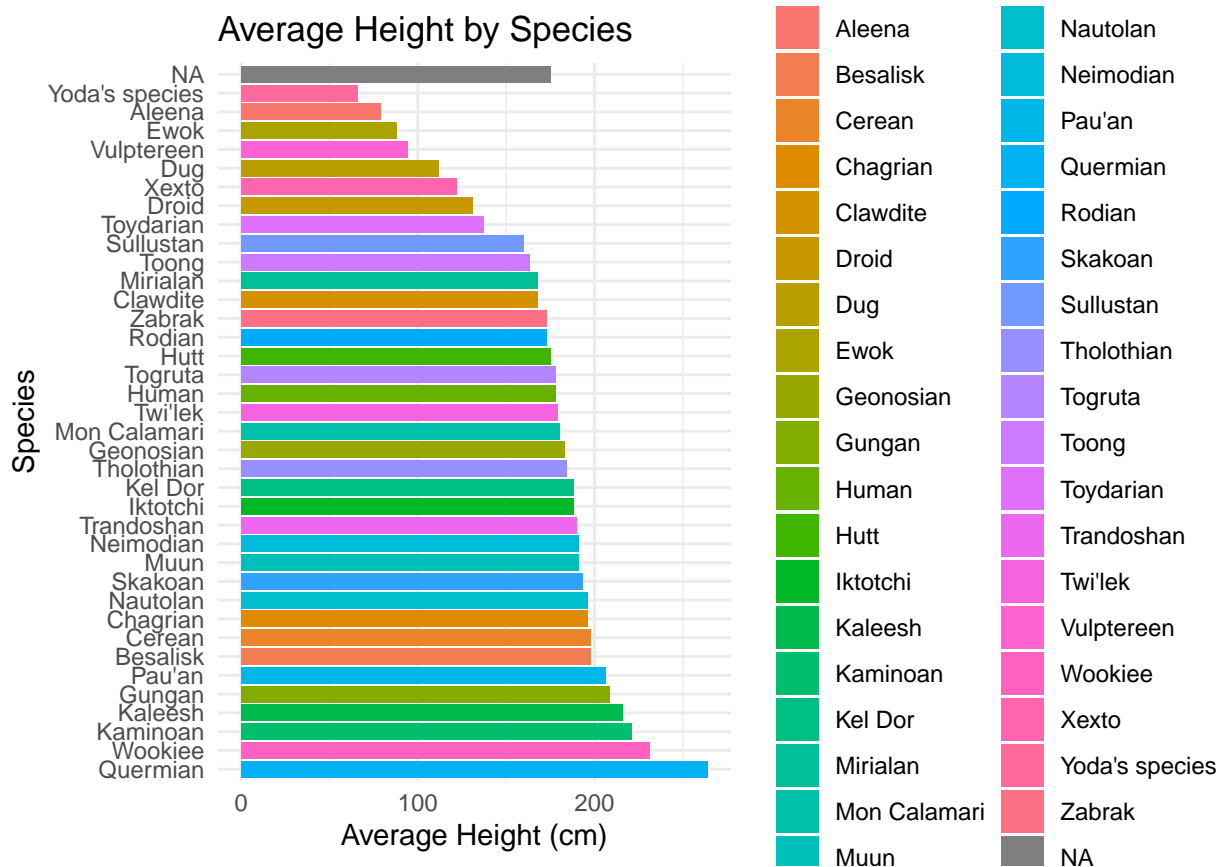
| | species | avg_height | avg_mass | count |
|---|----------|------------|----------|-------|
| | <chr> | <dbl> | <dbl> | <int> |
| 1 | Human | 178 | 81.3 | 35 |
| 2 | Droid | 131. | 69.8 | 6 |
| 3 | <NA> | 175 | 81 | 4 |
| 4 | Gungan | 209. | 74 | 3 |
| 5 | Kaminoan | 221 | 88 | 2 |
| 6 | Mirialan | 168 | 53.1 | 2 |

Plotting the average height

```

ggplot(species_summary, aes(x = reorder(species, -avg_height), y = avg_height, fill = species)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  labs(title = "Average Height by Species",
       x = "Species",
       y = "Average Height (cm)") +
  theme_minimal()

```



```
# Adding a new column that classifies characters based on height
```

```
starwars_classified <- starwars %>%
  mutate(height_category = ifelse(height < 180, "Short", "Tall"))
```

```
# Display the classified data
```

```
head(starwars_classified)
```

```
# A tibble: 6 x 15
```

```
  name    height  mass hair_color skin_color eye_color birth_year sex
  <chr>   <int> <dbl> <chr>      <chr>    <chr>      <dbl> <chr>
1 Luke ~    172    77 blond     fair     blue        19  male
2 C-3PO    167    75 <NA>      gold     yellow     112  none
3 R2-D2     96    32 <NA>      white, bl~ red        33  none
4 Darth~   202   136 none      white     yellow     41.9 male
5 Leia ~   150    49 brown     light    brown        19  fema~
6 Owen ~   178   120 brown, gr~ light     blue        52  male
# i 7 more variables: gender <chr>, homeworld <chr>, species <chr>,
#   films <list>, vehicles <list>, starships <list>,
#   height_category <chr>
```

```
# Plotting height Category distribution
ggplot(starwars_classified, aes(x = height_category, fill = height_category)) +
  geom_bar() +
  labs(title = "Height Category Distribution",
       x = "Height Category",
       y = "Count") +
  theme_minimal()
```



```
# Joining with another dataset (flights dataset from nycflights13)
data("flights")
data("airlines")

# Inner join flights with airlines on the common column "carrier"
flights_inner_join <- flights %>%
  inner_join(airlines, by = "carrier")

# Outer join flights with airlines on the common column "carrier"
flights_outer_join <- flights %>%
  full_join(airlines, by = "carrier")

# Display the joined data
head(flights_inner_join)
```

```
# A tibble: 6 x 20
  year month   day dep_time sched_dep_time dep_delay arr_time
  <int> <int> <int>   <int>         <int>       <dbl>   <int>
1  2013     1     1     517           515         2     830
2  2013     1     1     533           529         4     850
3  2013     1     1     542           540         2     923
4  2013     1     1     544           545        -1    1004
5  2013     1     1     554           600        -6     812
6  2013     1     1     554           558        -4     740
# i 13 more variables: sched_arr_time <int>, arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dtm>, name <chr>
```

```
head(flights_outer_join)
```

```
# A tibble: 6 x 20
  year month   day dep_time sched_dep_time dep_delay arr_time
  <int> <int> <int>   <int>         <int>       <dbl>   <int>
1  2013     1     1     517           515         2     830
2  2013     1     1     533           529         4     850
3  2013     1     1     542           540         2     923
4  2013     1     1     544           545        -1    1004
5  2013     1     1     554           600        -6     812
6  2013     1     1     554           558        -4     740
# i 13 more variables: sched_arr_time <int>, arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dtm>, name <chr>
```

```
# Calculating a 5 period rolling average of arrival delays and cumulative sum
```

```
flights_rolling <- flights %>%
  arrange(year, month, day) %>%
  mutate(
    arr_delay = ifelse(is.na(arr_delay), 0, arr_delay),
    rolling_avg_delay = zoo::rollmean(arr_delay, 5, fill = NA),
    cumulative_delay = cumsum(arr_delay)
  )
```

```
# Display the transformed data
```

```
head(flights_rolling)
```


```
# A tibble: 6 x 21
  year month   day dep_time sched_dep_time dep_delay arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>
1  2013     1     1     517             515           2     830
2  2013     1     1     533             529           4     850
3  2013     1     1     542             540           2     923
4  2013     1     1     544             545          -1    1004
5  2013     1     1     554             600          -6     812
6  2013     1     1     554             558          -4     740

# i 14 more variables: sched_arr_time <int>, arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dtm>, rolling_avg_delay <dbl>,
#   cumulative_delay <dbl>
```

Plotting the rolling average and cumulative delays

```
ggplot(flights_rolling, aes(x = day)) +
  geom_line(aes(y = rolling_avg_delay, color = "Rolling Average Delay")) +
  geom_line(aes(y = cumulative_delay / 1000, color = "Cumulative Delay (x1000)")) +
  labs(title = "Rolling Average and Cumulative Arrival Delays",
       x = "Day of the Month",
       y = "Delay (minutes)") +
  scale_color_manual(values = c("Rolling Average Delay" = "blue",
                                "Cumulative Delay (x1000)" = "red")) +
  theme_minimal()
```

Warning: Removed 4 rows containing missing values or values outside the scale range ('geom_line()').



Program - 6

Data Visualisation with ggplot2 and Customisations

Date of Execution - 2025-09-23


Objective - This program evaluates students' ability to create and customize complex data visualizations using the ggplot2 package.

```
# Load necessary libraries
library(ggplot2)
library(dplyr)
library(reshape2)

# Scatterplot with regression line and confidence intervals
data("mpg")


ggplot(mpg, aes(x = displ, y = hwy, color = class)) +
  geom_point(size = 3, alpha = 0.7) +
  geom_smooth(method = "lm", se = TRUE, linetype = "dashed") +
  labs(title = "Scatterplot of Engine Displacement vs Highway MPG",
       x = "Engine Displacement (liters)",
       y = "Highway MPG",
       color = "Vehicle Class") +
  theme_bw() +
  theme(plot.title = element_text(hjust = 0.5, size = 16, face = "bold"),
        axis.title = element_text(size = 14),
        legend.position = "bottom")
```

Scatterplot of Engine Displacement vs Highway MPG



```
# Multi-panel plot using Faceting
# Creating faceted scatter plots by vehicle class with enhanced aesthetics
ggplot(mpg, aes(x = displ, y = hwy)) +
  geom_point(color = "darkgreen", size = 2) +
  facet_wrap(~ class, ncol = 3) +
  labs(title = "Faceted Scatterplot by Vehicle Class",
       x = "Engine Displacement (liters)",
       y = "Highway MPG",
       color = "Drive Type") +
  theme_minimal() +
  theme(strip.text = element_text(size = 12, face = "italic"),
        plot.title = element_text(hjust = 0.5, size = 16))
```

Faceted Scatterplot by Vehicle Class



```
# Heatmap of correlatio matrix
```

```
data("diamonds")
```

```
# Calculate correlation matrix for numeric variables
```

```
cor_matrix <- cor(diamonds[sapply(diamonds, is.numeric)], use = "complete.obs")
```

```
#Convert to tidy format
```


```
cor_melt <- melt(cor_matrix)
```

```
# Create heatmap
```

```
ggplot(cor_melt, aes(Var1, Var2, fill = value)) +  
  geom_tile(color = "white") +  
  scale_fill_gradient2(low = "blue", high = "red", mid = "white",  
                        midpoint = 0, limit = c(-1, 1), space = "Lab",  
                        name = "Correlation") +  
  labs(title = "Heatmap of Correlation Matrix for Diamonds Dataset",  
        x = "Variables",  
        y = "Variables") +  
  theme_minimal() +  
  theme(axis.text.x = element_text(angle = 45, hjust = 1, size = 12),  
        axis.text.y = element_text(size = 12),
```

```
plot.title = element_text(hjust = 0.5, size = 16))
```


Heatmap of Correlation Matrix for Diamonds Dataset



Enhancing the scatterplot with annotations

```
ggplot(mpg, aes(x = displ, y = hwy, fill = class)) +
  geom_point(size = 3, shape = 21, alpha = 0.8) +
  theme_light() +
  scale_color_brewer(palette = "Set2") +
  labs(title = "Customised Scatter Plot",
       x = "Engine Displacement (liters)",
       y = "Highway MPG",
       color = "Vehicle Class") +
  theme(plot.title = element_text(face = "bold", size = 18),
        axis.title = element_text(size = 14),
        legend.background = element_rect(fill = "gray90"))
```

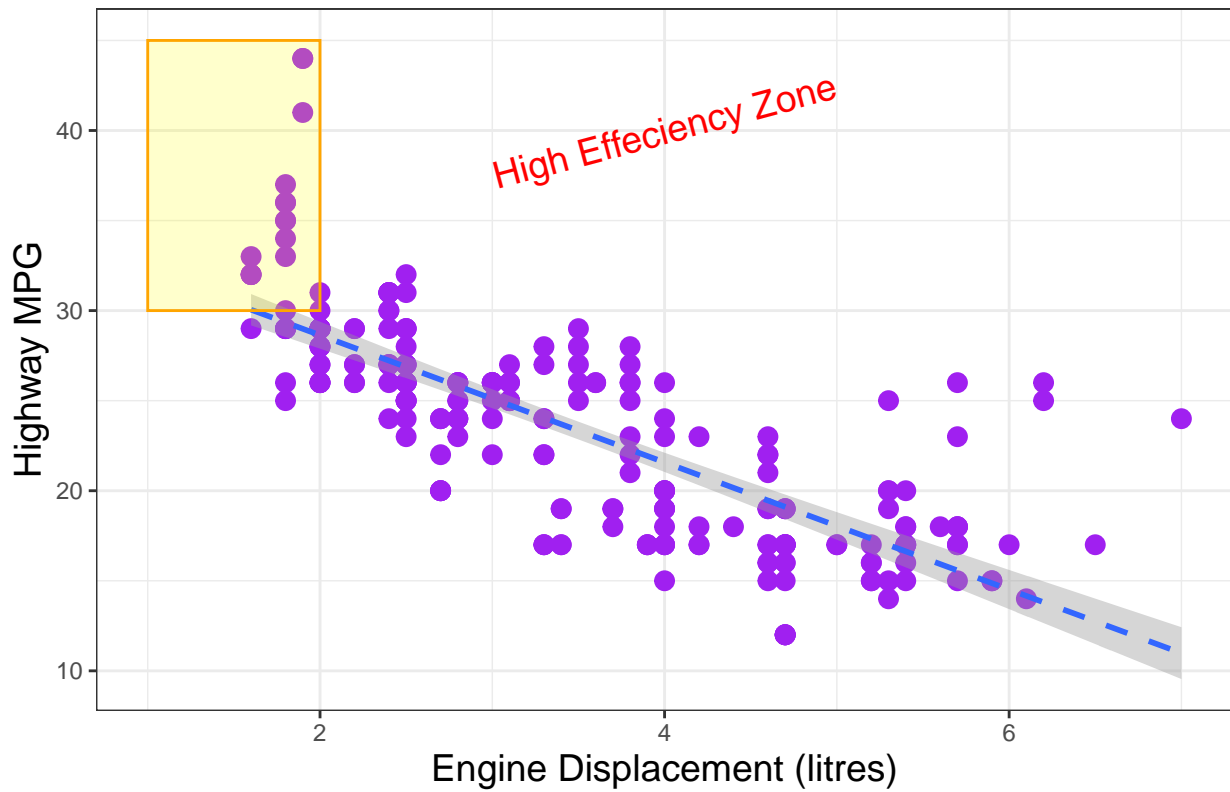
Customised Scatter Plot



```
# Annotate plots and save as image files
annotated_plot <- ggplot(mpg, aes(x = displ, y = hwy)) +
  geom_point(size = 3, color = "purple") +
  geom_smooth(method = "lm", se = TRUE, linetype = "dashed") +
  labs(title = "Annotate the Plot",
       x = "Engine Displacement (litres)",
       y = "Highway MPG",
       color = "Vehicle Class") +
  theme_bw() +
  theme(plot.title = element_text(hjust = 0.5, size = 16, face = "bold"),
        axis.title = element_text(size = 14),
        legend.position = "bottom") +
  annotate("text", x = 4, y = 40, label = "High Efficiency Zone",
         color = "red", size = 5, angle = 15) +
  annotate("rect", xmin = 1, xmax = 2, ymin = 30, ymax = 45,
         alpha = 0.2, fill = "yellow", color = "orange")

annotated_plot
```

Annotate the Plot



```
ggsave("annotated_scatterplot.png", plot = annotated_plot, width = 8, height = 6)
```

Program - 7

Linear and Multiple Regression Analysis with Interaction Terms

Date of Execution - 2025-10-14

Objective - This program focuses on regression modeling, interaction effects, and model diagnostics.

```
# Load necessary libraries
library(MASS)
library(ggplot2)
library(dplyr)
library(caret)
library(car)
library(pROC)
library(corrplot)

# Load the Boston Housing dataset
data("Boston")
head(Boston)
```

| | crim | zn | indus | chas | nox | rm | age | dis | rad | tax | ptratio |
|---|---------|----|-------|------|-------|-------|------|--------|-----|-----|---------|
| 1 | 0.00632 | 18 | 2.31 | 0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296 | 15.3 |
| 2 | 0.02731 | 0 | 7.07 | 0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | 17.8 |
| 3 | 0.02729 | 0 | 7.07 | 0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2 | 242 | 17.8 |
| 4 | 0.03237 | 0 | 2.18 | 0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | 18.7 |
| 5 | 0.06905 | 0 | 2.18 | 0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | 18.7 |
| 6 | 0.02985 | 0 | 2.18 | 0 | 0.458 | 6.430 | 58.7 | 6.0622 | 3 | 222 | 18.7 |

| | black | lstat | medv |
|---|--------|-------|------|
| 1 | 396.90 | 4.98 | 24.0 |
| 2 | 396.90 | 9.14 | 21.6 |
| 3 | 392.83 | 4.03 | 34.7 |
| 4 | 394.63 | 2.94 | 33.4 |
| 5 | 396.90 | 5.33 | 36.2 |
| 6 | 394.12 | 5.21 | 28.7 |

```
# q1. Preprocessing
```

```
# Check for missing values
```

```
sum(is.na(Boston))
```


[1] 0

```
# Summary Statistics
```

```
summary(Boston)
```

| crim | zn | indus | |
|------------------|----------------|----------------|--|
| Min. : 0.00632 | Min. : 0.00 | Min. : 0.46 | |
| 1st Qu.: 0.08205 | 1st Qu.: 0.00 | 1st Qu.: 5.19 | |
| Median : 0.25651 | Median : 0.00 | Median : 9.69 | |
| Mean : 3.61352 | Mean : 11.36 | Mean : 11.14 | |
| 3rd Qu.: 3.67708 | 3rd Qu.: 12.50 | 3rd Qu.: 18.10 | |
| Max. : 88.97620 | Max. : 100.00 | Max. : 27.74 | |

| chas | nox | rm | age |
|------------------|-----------------|----------------|----------------|
| Min. : 0.00000 | Min. : 0.3850 | Min. : 3.561 | Min. : 2.90 |
| 1st Qu.: 0.00000 | 1st Qu.: 0.4490 | 1st Qu.: 5.886 | 1st Qu.: 45.02 |
| Median : 0.00000 | Median : 0.5380 | Median : 6.208 | Median : 77.50 |
| Mean : 0.06917 | Mean : 0.5547 | Mean : 6.285 | Mean : 68.57 |
| 3rd Qu.: 0.00000 | 3rd Qu.: 0.6240 | 3rd Qu.: 6.623 | 3rd Qu.: 94.08 |
| Max. : 1.00000 | Max. : 0.8710 | Max. : 8.780 | Max. : 100.00 |


| dis | rad | tax | ptratio |
|----------------|-----------------|----------------|----------------|
| Min. : 1.130 | Min. : 1.000 | Min. : 187.0 | Min. : 12.60 |
| 1st Qu.: 2.100 | 1st Qu.: 4.000 | 1st Qu.: 279.0 | 1st Qu.: 17.40 |
| Median : 3.207 | Median : 5.000 | Median : 330.0 | Median : 19.05 |
| Mean : 3.795 | Mean : 9.549 | Mean : 408.2 | Mean : 18.46 |
| 3rd Qu.: 5.188 | 3rd Qu.: 24.000 | 3rd Qu.: 666.0 | 3rd Qu.: 20.20 |
| Max. : 12.127 | Max. : 24.000 | Max. : 711.0 | Max. : 22.00 |

| black | lstat | medv |
|-----------------|----------------|----------------|
| Min. : 0.32 | Min. : 1.73 | Min. : 5.00 |
| 1st Qu.: 375.38 | 1st Qu.: 6.95 | 1st Qu.: 17.02 |
| Median : 391.44 | Median : 11.36 | Median : 21.20 |
| Mean : 356.67 | Mean : 12.65 | Mean : 22.53 |
| 3rd Qu.: 396.23 | 3rd Qu.: 16.95 | 3rd Qu.: 25.00 |
| Max. : 396.90 | Max. : 37.97 | Max. : 50.00 |

```
# Check for outliers using boxplots
```

```
boxplot(Boston$medv, main = "Boxplot of Median Value (medv)")
```

Boxplot of Median Value (medv)



```
# Remove potential outliers (optional, based on domain knowledge)
```


```
Boston <- Boston %>% filter(medv < 50)
```

```
# 2. Feature Selection
```

```
# Calculate correlation matrix
```

```
corr_matrix <- cor(Boston)
```

```
corrplot(corr_matrix, method = "circle")
```



```
# High correlation observed between 'medv', 'lstat', and 'rm'
# We will use 'lstat' and 'rm' as predictors based on this analysis.
```

3. Simple Linear Regression Model

```
simple_model <- lm(medv ~ lstat, data = Boston)
summary(simple_model)
```

Call:

```
lm(formula = medv ~ lstat, data = Boston)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|--------|--------|-------|--------|
| -13.992 | -3.313 | -0.941 | 1.914 | 21.246 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|----------|------------|---------|------------|
| (Intercept) | 32.54041 | 0.48150 | 67.58 | <2e-16 *** |
| lstat | -0.84374 | 0.03268 | -25.82 | <2e-16 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.119 on 488 degrees of freedom
Multiple R-squared: 0.5774, Adjusted R-squared: 0.5765
F-statistic: 666.6 on 1 and 488 DF, p-value: < 2.2e-16

Interpretation

*# - The negative coefficient for 'lstat' suggests that higher 'lstat' values
(higher percentage of lower status population) are associated with lower 'medv'
(median home value)
- The p-value (<0.05) indicates that the relationship is statistically significant.*

4. Multiple Linear Regression

```
multiple_model <- lm(medv ~ lstat * rm, data = Boston)
summary(multiple_model)
```

Call:

```
lm(formula = medv ~ lstat * rm, data = Boston)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|----------|---------|---------|--------|---------|
| -21.3064 | -2.4982 | -0.3056 | 1.8635 | 18.4779 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|-----------|------------|---------|--------------|
| (Intercept) | -25.99970 | 3.07045 | -8.468 | 2.98e-16 *** |
| lstat | 1.97178 | 0.17761 | 11.102 | < 2e-16 *** |
| rm | 9.01216 | 0.46519 | 19.373 | < 2e-16 *** |
| lstat:rm | -0.43817 | 0.02976 | -14.723 | < 2e-16 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.845 on 486 degrees of freedom
Multiple R-squared: 0.7625, Adjusted R-squared: 0.761
F-statistic: 520 on 3 and 486 DF, p-value: < 2.2e-16

Interpretation:

*# - Significant coefficients for 'lstat', 'rm', and the interaction term ('lstat:rm')
- Indicates that the relationship between 'lstat' and 'medv' depends on the value of 'rm'
- The adjusted R² has improved, suggesting better fit when compared to simple model*

5. Model Performance Evaluation

```
adjusted_R2 <- summary(multiple_model)$adj.r.squared
AIC_Value <- AIC(multiple_model)
BIC_Value <- BIC(multiple_model)

cat("Adjusted R^2: ", adjusted_R2 , "\n")
```

Adjusted R²: 0.7609872

```
cat("AIC :", AIC_Value , "\n")
```

AIC : 2716.448


```
cat("BIC :", BIC_Value , "\n")
```

BIC : 2737.42


#6. Model Diagnostics : Residual Analysis

Residual v/s Fitted plot

```
plot(multiple_model, which = 1, main = "Residuals v/s Fitted Plot")
```



```
# Q-Q plot for checking normality of residuals
plot(multiple_model, which = 2, main = "Normal Q-Q Plot")
```



```
# Interpretation
# - The residuals show a random scatter around zero in the Residuals v/s Fitted Plot
# - The Q-Q plot follows a straight line if the residuals are normally distributed

# 7. Cross Validation Model Accuracy
set.seed(123)
train_control <- trainControl(method = "cv", number = 10)
cv_model <- train(medv ~ lstat*rm, data = Boston,
                  method = "lm",
                  trControl = train_control)

# Results
print(cv_model)
```

Linear Regression

490 samples

2 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 441, 441, 442, 441, 441, 440, ...

Resampling results:

| RMSE | Rsquared | MAE |
|----------|-----------|----------|
| 3.855714 | 0.7597212 | 2.868657 |

Tuning parameter 'intercept' was held constant at a value of TRUE

Interpretation:

- Cross Validation RMSE provides a estimate of prediction error

- Lower RMSE indicates better model performance

8. ROC Curve Analysis (Classification Approach)

Convert 'medv' to a binary classification problem: High (≥ 25) or Low (< 25)

```
Boston$medv_class <- ifelse(Boston$medv >= 25, 1, 0)
```

Fit a logistic regression model for classification

```
logistic_model <- glm(medv_class ~ lstat*rm, data = Boston, family = "binomial")
```

```
summary(logistic_model)
```

Call:

```
glm(formula = medv_class ~ lstat * rm, family = "binomial", data = Boston)
```

Coefficients:

| | Estimate | Std. Error | z value | Pr(> z) |
|-------------|-----------|------------|---------|------------|
| (Intercept) | -19.80238 | 7.45022 | -2.658 | 0.00786 ** |
| lstat | -0.02301 | 0.75334 | -0.031 | 0.97563 |
| rm | 3.29921 | 1.12727 | 2.927 | 0.00343 ** |
| lstat:rm | -0.03989 | 0.11492 | -0.347 | 0.72851 |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 536.34 on 489 degrees of freedom


Residual deviance: 254.70 on 486 degrees of freedom

AIC: 262.7

Number of Fisher Scoring iterations: 8

```
# Predict probabilities and compute ROC Curve
pred_probs <- predict(logistic_model, type = "response")
roc_curve <- roc(Boston$medv_class, pred_probs)

# Plot ROC Curves
plot(roc_curve, main = "ROC Curve for Logistic Regression Model", col = "blue")
abline(a=0, b=1, lty = 2, col = "red")
```



```
cat("AUC: ", auc(roc_curve), "\n")
```

AUC: 0.9392864

```
# Interpretation
# - The ROC Curve evaluates the trade-off between sensitivity and specificity
# - The Area Under the Curve (AUC) indicates the model's discriminatory ability
# (AUC closer to 1 is better)
```


Program - 8

K-Means Clustering and PCA for Dimensionality Reduction

Date of Execution - 2025-10-28

Objective - This program tests the student's knowledge of clustering techniques and dimensionality reduction through PCA.

```
# Load required libraries
library(rattle)      # For Wine dataset
library(ggplot2)     # For visualization
library(cluster)     # For silhouette scores
library(factoextra)  # For PCA and clustering visualization
library(dplyr)       # Often useful for data manipulation

# Normalize function (Min-Max Scaling)
normalize <- function(data) {
  return((data - min(data)) / (max(data) - min(data)))
}

# -----
# Analysis for WINE Dataset
# -----

# Step 1: Load Wine dataset and normalize
data(wine)
wine_data <- wine[, -1] # Remove the class label
wine_norm <- as.data.frame(lapply(wine_data, normalize))

# Step 2: Apply PCA
wine_pca <- prcomp(wine_norm, scale. = TRUE)
summary(wine_pca)
```

Importance of components:

| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 |
|------------------------|-------|--------|--------|---------|---------|---------|
| Standard deviation | 2.169 | 1.5802 | 1.2025 | 0.95863 | 0.92370 | 0.80103 |
| Proportion of Variance | 0.362 | 0.1921 | 0.1112 | 0.07069 | 0.06563 | 0.04936 |
| Cumulative Proportion | 0.362 | 0.5541 | 0.6653 | 0.73599 | 0.80162 | 0.85098 |
| | PC7 | PC8 | PC9 | PC10 | PC11 | PC12 |

```
Standard deviation      0.74231 0.59034 0.53748 0.5009 0.47517 0.41082
Proportion of Variance 0.04239 0.02681 0.02222 0.0193 0.01737 0.01298
Cumulative Proportion  0.89337 0.92018 0.94240 0.9617 0.97907 0.99205
```

PC13

```
Standard deviation      0.32152
Proportion of Variance 0.00795
Cumulative Proportion  1.00000
```


```
# Reduce to top 2 principal components
```

```
wine_pca_data <- as.data.frame(wine_pca$x[, 1:2])
```

```
# Step 3: Determine the optimal number of clusters (Elbow method)
```

```
elbow_wine <- fviz_nbclust(wine_pca_data, kmeans, method = "wss")
```


```
print(elbow_wine)
```



```
# Step 4: Silhouette analysis
```

```
silhouette_wine <- fviz_nbclust(wine_pca_data, kmeans, method = "silhouette")
```

```
print(silhouette_wine)
```




```
# Step 5: Apply K-means clustering (using centers=3 based on known structure/analysis)
set.seed(123)
wine_kmeans <- kmeans(wine_pca_data, centers = 3, nstart = 25)

# Step 6: Visualize clusters
wine_pca_data$cluster <- as.factor(wine_kmeans$cluster)

p1 <- ggplot(wine_pca_data, aes(x = PC1, y = PC2, color = cluster)) +
  geom_point(size = 3) +
  labs(title = "K-Means Clustering on Wine Dataset")
print(p1)
```

K-Means Clustering on Wine Dataset



Step 7: Interpret results

```
cat("Wine Dataset Clustering Results:\n")
```

Wine Dataset Clustering Results:

```
cat("Cluster Sizes: ", wine_kmeans$size, "\n")
```

Cluster Sizes: 64 65 49

```
# -----
```

Step 8: Analysis for Breast Cancer Wisconsin Dataset

```
# -----
```

Load dataset from UCI repository (Ensure you have internet connection)

```
bc_data <- read.csv("https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/")
```

```
bc_features <- bc_data[, -c(1, 2)] # Exclude ID and Class columns
```

```
bc_norm <- as.data.frame(lapply(bc_features, normalize))
```

Apply PCA

```
bc_pca <- prcomp(bc_norm, scale. = TRUE)
```

```
summary(bc_pca)
```

Importance of components:

| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 |
|------------------------|---------|---------|---------|---------|---------|---------|
| Standard deviation | 3.6444 | 2.3857 | 1.67867 | 1.40735 | 1.28403 | 1.09880 |
| Proportion of Variance | 0.4427 | 0.1897 | 0.09393 | 0.06602 | 0.05496 | 0.04025 |
| Cumulative Proportion | 0.4427 | 0.6324 | 0.72636 | 0.79239 | 0.84734 | 0.88759 |
| | PC7 | PC8 | PC9 | PC10 | PC11 | PC12 |
| Standard deviation | 0.82172 | 0.69037 | 0.6457 | 0.59219 | 0.5421 | 0.51104 |
| Proportion of Variance | 0.02251 | 0.01589 | 0.0139 | 0.01169 | 0.0098 | 0.00871 |
| Cumulative Proportion | 0.91010 | 0.92598 | 0.9399 | 0.95157 | 0.9614 | 0.97007 |
| | PC13 | PC14 | PC15 | PC16 | PC17 | |
| Standard deviation | 0.49128 | 0.39624 | 0.30681 | 0.28260 | 0.24372 | |
| Proportion of Variance | 0.00805 | 0.00523 | 0.00314 | 0.00266 | 0.00198 | |
| Cumulative Proportion | 0.97812 | 0.98335 | 0.98649 | 0.98915 | 0.99113 | |
| | PC18 | PC19 | PC20 | PC21 | PC22 | PC23 |
| Standard deviation | 0.22939 | 0.22244 | 0.17652 | 0.1731 | 0.16565 | 0.15602 |
| Proportion of Variance | 0.00175 | 0.00165 | 0.00104 | 0.0010 | 0.00091 | 0.00081 |
| Cumulative Proportion | 0.99288 | 0.99453 | 0.99557 | 0.9966 | 0.99749 | 0.99830 |
| | PC24 | PC25 | PC26 | PC27 | PC28 | PC29 |
| Standard deviation | 0.1344 | 0.12442 | 0.09043 | 0.08307 | 0.03987 | 0.02736 |
| Proportion of Variance | 0.0006 | 0.00052 | 0.00027 | 0.00023 | 0.00005 | 0.00002 |
| Cumulative Proportion | 0.9989 | 0.99942 | 0.99969 | 0.99992 | 0.99997 | 1.00000 |
| | PC30 | | | | | |
| Standard deviation | 0.01153 | | | | | |
| Proportion of Variance | 0.00000 | | | | | |
| Cumulative Proportion | 1.00000 | | | | | |

Reduce to top 2 principal components


```
bc_pca_data <- as.data.frame(bc_pca$x[, 1:2])
```

Optimal number of clusters (Elbow method)

```
elbow_bc <- fviz_nbclust(bc_pca_data, kmeans, method = "wss")
```

```
print(elbow_bc)
```

Optimal number of clusters




```
# Optimal number of clusters (Silhouette analysis)
```

```
silhouette_bc <- fviz_nbclust(bc_pca_data, kmeans, method = "silhouette")
```

```
print(silhouette_bc)
```

Optimal number of clusters




```

# Apply K-means clustering (using centers=2 based on binary classification/analysis)
set.seed(123)
bc_kmeans <- kmeans(bc_pca_data, centers = 2, nstart = 25)

# Add cluster assignments to PCA data
bc_pca_data$cluster <- as.factor(bc_kmeans$cluster)

# Visualize clusters
p2 <- ggplot(bc_pca_data, aes(x = PC1, y = PC2, color = cluster)) +
  geom_point(size = 3) +
  labs(title = "K-Means Clustering on Breast Cancer Dataset")
print(p2)

```



```

# Interpret results
cat("Breast Cancer Dataset Clustering Results:\n")

```

Breast Cancer Dataset Clustering Results:

```
cat("Cluster Sizes: ", bc_kmeans$size, "\n")
```

Cluster Sizes: 191 378

Program - 9

Time Series Analysis using ARIMA and Seasonal Decomposition

Date of Execution - 2025-10-28

Objective - This program evaluates students' skills in time series analysis, model fitting, and forecasting.

```
# Load required libraries for time series analysis and modeling
library(forecast)
library(ggplot2)
library(TSA)
library(tseries)

# Function to perform Exploratory Data Analysis (EDA) on the time series data
perform_eda <- function(ts_data, dataset_name) {
  cat(" Exploratory Data Analysis for ", dataset_name, "\n")
  print(summary(ts_data)) # Print summary of the dataset
  plot(ts_data, main = paste(dataset_name, " Time Series "), ylab = " Values ", xlab = "Time ")
  cat("ACF and PACF plots :\n")
  acf(ts_data, main = paste("ACF of", dataset_name)) # Autocorrelation plot
  pacf(ts_data, main = paste(" PACF of", dataset_name)) # Partial autocorrelation plot
}

# Function to decompose the time series into trend , seasonal , and residual components
decompose_ts <- function(ts_data, dataset_name) {
  cat(" Decomposing the time series for ", dataset_name, "\n")
  decomposition <- decompose(ts_data) # Decompose the time series
  plot(decomposition) # Plot the decomposition
  return(decomposition) # Return the decomposition result
}

# Function to fit an ARIMA model to the time series data
fit_arima <- function(ts_data, dataset_name) {
  cat(" Fitting ARIMA model for ", dataset_name, "\n")
  adf_test <- adf.test(ts_data, alternative = "stationary") # ADF test for stationarity
  cat("ADF Test p- value :", adf_test$p.value, "\n")
  # If p- value > 0.05 , data is non - stationary , so we difference the data
  if (adf_test$p.value > 0.05) {
    ts_data <- diff(ts_data) # Difference the data to make it stationary
  }
}
```



```

    plot(ts_data, main = paste(dataset_name, " Differenced Time Series "))
  }
  auto_model <- auto.arima(ts_data, seasonal = FALSE) # Fit ARIMA model (non - seasonal)
  print(summary(auto_model)) # Print ARIMA model summary
  forecast_result <- forecast(auto_model, h = 12) # Forecast next 12 periods
  plot(forecast_result, main = paste(dataset_name, " ARIMA Forecast ")) # Plot ARIMA forecast
  return(auto_model) # Return the fitted ARIMA model
}

# Function to fit a Seasonal ARIMA ( SARIMA ) model to the time series data
fit_sarima <- function(ts_data, dataset_name) {
  cat(" Fitting SARIMA model for ", dataset_name, "\n")
  auto_sarima <- auto.arima(ts_data, seasonal = TRUE) # Fit SARIMA model ( seasonal )
  print(summary(auto_sarima)) # Print SARIMA model summary
  sarima_forecast <- forecast(auto_sarima, h = 12) # Forecast next 12 periods
  plot(sarima_forecast, main = paste(dataset_name, " SARIMA Forecast ")) # Plot SARIMA forecast
  return(auto_sarima) # Return the fitted SARIMA model
}

# Function to compare ARIMA and SARIMA models by evaluating forecast accuracy
compare_models <- function(arima_model, sarima_model, ts_data) {
  cat(" Comparing ARIMA and SARIMA models :\n")
  h <- min(12, length(ts_data)) # Forecast horizon of 12 or adjusted based on dataset length
  arima_forecast <- forecast(arima_model, h = h) # ARIMA forecast
  sarima_forecast <- forecast(sarima_model, h = h) # SARIMA forecast
  actual_values <- ts_data[(length(ts_data) - h + 1): length(ts_data)] # Comparison

  # Calculate accuracy of both models
  arima_accuracy <- accuracy(arima_forecast$mean, actual_values)
  sarima_accuracy <- accuracy(sarima_forecast$mean, actual_values)
  cat(" ARIMA Forecast Accuracy :\n", arima_accuracy) # Print ARIMA accuracy
  cat(" SARIMA Forecast Accuracy :\n", sarima_accuracy) # Print SARIMA accuracy
}

# Function to visualize the comparison of ARIMA and SARIMA forecast performance
plot_forecast_comparison <- function(actual_values, arima_forecast, sarima_forecast, time_points) {
  arima_rmse <- sqrt(mean((arima_forecast - actual_values)^2)) # Calculate RMSE for ARIMA
  sarima_rmse <- sqrt(mean((sarima_forecast - actual_values)^2)) # Calculate RMSE for SARIMA

  # Color coding for better and worse RMSE

```

```

better_color <- ifelse(arima_rmse < sarima_rmse, "green", "red")
worse_color <- ifelse(arima_rmse < sarima_rmse, "red", "green")

# Plot actual values and forecasts
plot(time_points, actual_values, type = "o", col = "blue", pch = 16, lty = 1, xlab = " Time ",
      ylab = " Values ", main = " Forecast Comparison ")
lines(time_points, arima_forecast, col = better_color, lty = 2, lwd = 2) # ARIMA forecast line
lines(time_points, sarima_forecast, col = worse_color, lty = 3, lwd = 2) # SARIMA forecast line

# Add a legend to the plot
legend("topright", legend = c("Actual Values", paste("ARIMA (RMSE =", round(arima_rmse, 2), ")"),
                           paste("SARIMA (RMSE =", round(sarima_rmse, 2), ")")),
      col = c("blue", better_color, worse_color), lty = c(1, 2, 3), lwd = c(1, 2, 2),
      pch = c(16, NA, NA))
}

# AirPassengers Dataset Analysis
data("AirPassengers")
air_data <- AirPassengers
cat("\n- - - AirPassengers Dataset - - -\n")

```


- - - AirPassengers Dataset - - -

```
perform_eda(air_data, "AirPassengers")
```

Exploratory Data Analysis for AirPassengers


| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|-------|---------|--------|-------|---------|-------|
| 104.0 | 180.0 | 265.5 | 280.3 | 360.5 | 622.0 |

AirPassengers Time Series




ACF and PACF plots :

ACF of AirPassengers




PACF of AirPassengers



```
decompose_ts(air_data, "AirPassengers")
```

Decomposing the time series for AirPassengers

Decomposition of additive time series



\$x

| | Jan | Feb | Mar | Apr | May | Jun | Jul | Aug | Sep | Oct | Nov | Dec |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 1949 | 112 | 118 | 132 | 129 | 121 | 135 | 148 | 148 | 136 | 119 | 104 | 118 |
| 1950 | 115 | 126 | 141 | 135 | 125 | 149 | 170 | 170 | 158 | 133 | 114 | 140 |
| 1951 | 145 | 150 | 178 | 163 | 172 | 178 | 199 | 199 | 184 | 162 | 146 | 166 |
| 1952 | 171 | 180 | 193 | 181 | 183 | 218 | 230 | 242 | 209 | 191 | 172 | 194 |
| 1953 | 196 | 196 | 236 | 235 | 229 | 243 | 264 | 272 | 237 | 211 | 180 | 201 |
| 1954 | 204 | 188 | 235 | 227 | 234 | 264 | 302 | 293 | 259 | 229 | 203 | 229 |
| 1955 | 242 | 233 | 267 | 269 | 270 | 315 | 364 | 347 | 312 | 274 | 237 | 278 |
| 1956 | 284 | 277 | 317 | 313 | 318 | 374 | 413 | 405 | 355 | 306 | 271 | 306 |
| 1957 | 315 | 301 | 356 | 348 | 355 | 422 | 465 | 467 | 404 | 347 | 305 | 336 |
| 1958 | 340 | 318 | 362 | 348 | 363 | 435 | 491 | 505 | 404 | 359 | 310 | 337 |
| 1959 | 360 | 342 | 406 | 396 | 420 | 472 | 548 | 559 | 463 | 407 | 362 | 405 |
| 1960 | 417 | 391 | 419 | 461 | 472 | 535 | 622 | 606 | 508 | 461 | 390 | 432 |

\$seasonal

| | Jan | Feb | Mar | Apr | May |
|------|------------|------------|-----------|-----------|------------|
| 1949 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1950 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1951 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1952 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1953 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1954 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1955 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1956 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1957 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1958 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1959 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| 1960 | -24.748737 | -36.188131 | -2.241162 | -8.036616 | -4.506313 |
| | Jun | Jul | Aug | Sep | Oct |
| 1949 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1950 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1951 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1952 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1953 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1954 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1955 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1956 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1957 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1958 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |
| 1959 | 35.402778 | 63.830808 | 62.823232 | 16.520202 | -20.642677 |

1960 35.402778 63.830808 62.823232 16.520202 -20.642677

Nov Dec

1949 -53.593434 -28.619949

1950 -53.593434 -28.619949

1951 -53.593434 -28.619949

1952 -53.593434 -28.619949

1953 -53.593434 -28.619949

1954 -53.593434 -28.619949

1955 -53.593434 -28.619949

1956 -53.593434 -28.619949

1957 -53.593434 -28.619949

1958 -53.593434 -28.619949

1959 -53.593434 -28.619949

1960 -53.593434 -28.619949

\$trend

Jan Feb Mar Apr May Jun Jul

1949 NA NA NA NA NA NA 126.7917

1950 131.2500 133.0833 134.9167 136.4167 137.4167 138.7500 140.9167

1951 157.1250 159.5417 161.8333 164.1250 166.6667 169.0833 171.2500

1952 183.1250 186.2083 189.0417 191.2917 193.5833 195.8333 198.0417

1953 215.8333 218.5000 220.9167 222.9167 224.0833 224.7083 225.3333

1954 228.0000 230.4583 232.2500 233.9167 235.6250 237.7500 240.5000

1955 261.8333 266.6667 271.1250 275.2083 278.5000 281.9583 285.7500

1956 309.9583 314.4167 318.6250 321.7500 324.5000 327.0833 329.5417

1957 348.2500 353.0000 357.6250 361.3750 364.5000 367.1667 369.4583

1958 375.2500 377.9167 379.5000 380.0000 380.7083 380.9583 381.8333

1959 402.5417 407.1667 411.8750 416.3333 420.5000 425.5000 430.7083

1960 456.3333 461.3750 465.2083 469.3333 472.7500 475.0417 NA

Aug Sep Oct Nov Dec

1949 127.2500 127.9583 128.5833 129.0000 129.7500

1950 143.1667 145.7083 148.4167 151.5417 154.7083

1951 173.5833 175.4583 176.8333 178.0417 180.1667

1952 199.7500 202.2083 206.2500 210.4167 213.3750

1953 225.3333 224.9583 224.5833 224.4583 225.5417

1954 243.9583 247.1667 250.2500 253.5000 257.1250

1955 289.3333 293.2500 297.1667 301.0000 305.4583

1956 331.8333 334.4583 337.5417 340.5417 344.0833

1957 371.2083 372.1667 372.4167 372.7500 373.6250

1958 383.6667 386.5000 390.3333 394.7083 398.6250

1959 435.1250 437.7083 440.9583 445.8333 450.6250

| | | | | | |
|----------|-------------|-------------|-------------|-------------|-------------|
| 1960 | NA | NA | NA | NA | NA |
| \$random | | | | | |
| | Jan | Feb | Mar | Apr | May |
| 1949 | NA | NA | NA | NA | NA |
| 1950 | 8.4987374 | 29.1047980 | 8.3244949 | 6.6199495 | -7.9103535 |
| 1951 | 12.6237374 | 26.6464646 | 18.4078283 | 6.9116162 | 9.8396465 |
| 1952 | 12.6237374 | 29.9797980 | 6.1994949 | -2.2550505 | -6.0770202 |
| 1953 | 4.9154040 | 13.6881313 | 17.3244949 | 20.1199495 | 9.4229798 |
| 1954 | 0.7487374 | -6.2702020 | 4.9911616 | 1.1199495 | 2.8813131 |
| 1955 | 4.9154040 | 2.5214646 | -1.8838384 | 1.8282828 | -3.9936869 |
| 1956 | -1.2095960 | -1.2285354 | 0.6161616 | -0.7133838 | -1.9936869 |
| 1957 | -8.5012626 | -15.8118687 | 0.6161616 | -5.3383838 | -4.9936869 |
| 1958 | -10.5012626 | -23.7285354 | -15.2588384 | -23.9633838 | -13.2020202 |
| 1959 | -17.7929293 | -28.9785354 | -3.6338384 | -12.2967172 | 4.0063131 |
| 1960 | -14.5845960 | -34.1868687 | -43.9671717 | -0.2967172 | 3.7563131 |
| | Jun | Jul | Aug | Sep | Oct |
| 1949 | NA | -42.6224747 | -42.0732323 | -8.4785354 | 11.0593434 |
| 1950 | -25.1527778 | -34.7474747 | -35.9898990 | -4.2285354 | 5.2260101 |
| 1951 | -26.4861111 | -36.0808081 | -37.4065657 | -7.9785354 | 5.8093434 |
| 1952 | -13.2361111 | -31.8724747 | -20.5732323 | -9.7285354 | 5.3926768 |
| 1953 | -17.1111111 | -25.1641414 | -16.1565657 | -4.4785354 | 7.0593434 |
| 1954 | -9.1527778 | -2.3308081 | -13.7815657 | -4.6868687 | -0.6073232 |
| 1955 | -2.3611111 | 14.4191919 | -5.1565657 | 2.2297980 | -2.5239899 |
| 1956 | 11.5138889 | 19.6275253 | 10.3434343 | 4.0214646 | -10.8989899 |
| 1957 | 19.4305556 | 31.7108586 | 32.9684343 | 15.3131313 | -4.7739899 |
| 1958 | 18.6388889 | 45.3358586 | 58.5101010 | 0.9797980 | -10.6906566 |
| 1959 | 11.0972222 | 53.4608586 | 61.0517677 | 8.7714646 | -13.3156566 |
| 1960 | 24.5555556 | NA | NA | NA | NA |
| | Nov | Dec | | | |
| 1949 | 28.5934343 | 16.8699495 | | | |
| 1950 | 16.0517677 | 13.9116162 | | | |
| 1951 | 21.5517677 | 14.4532828 | | | |
| 1952 | 15.1767677 | 9.2449495 | | | |
| 1953 | 9.1351010 | 4.0782828 | | | |
| 1954 | 3.0934343 | 0.4949495 | | | |
| 1955 | -10.4065657 | 1.1616162 | | | |
| 1956 | -15.9482323 | -9.4633838 | | | |
| 1957 | -14.1565657 | -9.0050505 | | | |
| 1958 | -31.1148990 | -33.0050505 | | | |
| 1959 | -30.2398990 | -17.0050505 | | | |

1960 NA NA

\$figure

```
[1] -24.748737 -36.188131 -2.241162 -8.036616 -4.506313 35.402778  
[7] 63.830808 62.823232 16.520202 -20.642677 -53.593434 -28.619949
```

\$type

```
[1] "additive"
```

attr(,"class")

```
[1] "decomposed.ts"
```

```
arima_air <- fit_arima(air_data, "AirPassengers")
```

Fitting ARIMA model for AirPassengers

Warning in adf.test(ts_data, alternative = "stationary"): p-value
smaller than printed p-value

ADF Test p- value : 0.01

Series: ts_data

ARIMA(4,1,2) with drift

Coefficients:

| | ar1 | ar2 | ar3 | ar4 | ma1 | ma2 | drift |
|------|--------|--------|---------|---------|---------|---------|--------|
| | 0.2243 | 0.3689 | -0.2567 | -0.2391 | -0.0971 | -0.8519 | 2.6809 |
| s.e. | 0.1047 | 0.1147 | 0.0985 | 0.0919 | 0.0866 | 0.0877 | 0.1711 |

sigma^2 = 706.3: log likelihood = -670.07

AIC=1356.15 AICc=1357.22 BIC=1379.85


Training set error measures:

| | ME | RMSE | MAE | MPE | MAPE | MASE |
|--------------|-----------|----------|----------|-----------|----------|-----------|
| Training set | -1.228696 | 25.82793 | 20.59211 | -1.665245 | 7.476447 | 0.6428946 |

ACF1

Training set 0.0009861078

AirPassengers ARIMA Forecast



```
sarima_air <- fit_sarima(air_data, "AirPassengers")
```

Fitting SARIMA model for AirPassengers

Series: ts_data

ARIMA(2,1,1)(0,1,0)[12]

Coefficients:

| | ar1 | ar2 | ma1 |
|------|--------|--------|---------|
| | 0.5960 | 0.2143 | -0.9819 |
| s.e. | 0.0888 | 0.0880 | 0.0292 |

sigma² = 132.3: log likelihood = -504.92

AIC=1017.85 AICc=1018.17 BIC=1029.35


Training set error measures:

| | ME | RMSE | MAE | MPE | MAPE | MASE |
|--------------|--------|----------|---------|----------|----------|----------|
| Training set | 1.3423 | 10.84619 | 7.86754 | 0.420698 | 2.800458 | 0.245628 |

ACF1

Training set -0.00124847

AirPassengers SARIMA Forecast



```
compare_models(arima_air, sarima_air, air_data)
```

Comparing ARIMA and SARIMA models :

ARIMA Forecast Accuracy :

-23.06901 83.58979 74.65149 -7.376947 16.02863 SARIMA Forecast Accuracy :

-31.66945 31.70666 31.66945 -6.784721 6.784721

```
# Forecasting and plot comparison for AirPassengers dataset
```

```
h_air <- 12 # Define forecast horizon for AirPassengers dataset (12 months ahead)
```

```
# Extract the actual values for the last 12 months of the AirPassengers data
```

```
air_actual_values <- air_data[(length(air_data) - h_air + 1): length(air_data)]
```

```
# Generate ARIMA forecast for the next 12 months
```

```
arima_air_forecast <- forecast(arima_air, h = h_air)$mean
```

```
# Generate SARIMA forecast for the next 12 months
```

```
sarima_air_forecast <- forecast(sarima_air, h = h_air)$mean
```


```
# Extract the time points for the last 12 months
```

```
time_points_air <- time(air_data)[(length(air_data) - h_air + 1): length(air_data)]
```

```
# Plot and compare the forecasts from ARIMA and SARIMA models against the actual values
```

```
plot_forecast_comparison(air_actual_values, arima_air_forecast, sarima_air_forecast,  
                          time_points_air)
```

Forecast Comparison



```
# Monthly Milk Production Dataset Analysis
```

```
data(milk) # Load the Monthly Milk Production dataset
```

```
milk_data <- milk # Assign the dataset to a variable
```

```
cat("\n- - Monthly Milk Production Dataset - - \n")
```

```
- - - Monthly Milk Production Dataset - - -
```

```
# Perform Exploratory Data Analysis (EDA) for the Milk Production dataset
```

```
perform_eda(milk_data, "Monthly Milk Production")
```

```
Exploratory Data Analysis for  Monthly Milk Production
```

```
    milk
```

```
Min.   :1236
```

```
1st Qu.:1420
```


```
Median :1504
```

```
Mean   :1504
```

```
3rd Qu.:1588
```


```
Max.   :1760
```

Monthly Milk Production Time Series




ACF and PACF plots :

ACF of Monthly Milk Production




PACF of Monthly Milk Production



```
# Decompose the Milk Production time series into trend , seasonal , and residual components  
decompose_ts(milk_data, "Monthly Milk Production")
```

Decomposing the time series for Monthly Milk Production

Decomposition of additive time series



\$x

| | Jan | Feb | Mar | Apr | May | Jun | Jul | Aug | Sep | Oct | Nov | Dec |
|------|------|------|------|------|------|------|------|------|------|------|------|------|
| 1994 | 1343 | 1236 | 1401 | 1396 | 1457 | 1388 | 1389 | 1369 | 1318 | 1354 | 1312 | 1370 |
| 1995 | 1404 | 1295 | 1453 | 1427 | 1484 | 1421 | 1414 | 1375 | 1331 | 1364 | 1320 | 1380 |
| 1996 | 1415 | 1348 | 1469 | 1441 | 1479 | 1398 | 1400 | 1382 | 1342 | 1391 | 1350 | 1418 |
| 1997 | 1433 | 1328 | 1500 | 1474 | 1529 | 1471 | 1473 | 1446 | 1377 | 1416 | 1369 | 1438 |
| 1998 | 1466 | 1347 | 1515 | 1501 | 1556 | 1477 | 1468 | 1443 | 1386 | 1446 | 1407 | 1489 |
| 1999 | 1518 | 1404 | 1585 | 1554 | 1610 | 1516 | 1498 | 1487 | 1445 | 1491 | 1459 | 1538 |
| 2000 | 1579 | 1506 | 1632 | 1593 | 1636 | 1547 | 1561 | 1525 | 1464 | 1511 | 1459 | 1519 |
| 2001 | 1549 | 1431 | 1599 | 1571 | 1632 | 1555 | 1552 | 1520 | 1472 | 1522 | 1485 | 1549 |
| 2002 | 1591 | 1472 | 1654 | 1621 | 1678 | 1587 | 1578 | 1570 | 1497 | 1539 | 1496 | 1575 |
| 2003 | 1615 | 1489 | 1666 | 1627 | 1671 | 1596 | 1597 | 1571 | 1511 | 1561 | 1517 | 1596 |
| 2004 | 1624 | 1531 | 1661 | 1636 | 1692 | 1607 | 1623 | 1601 | 1533 | 1583 | 1531 | 1610 |
| 2005 | 1643 | 1522 | 1707 | 1690 | 1760 | 1690 | 1683 | 1671 | 1599 | 1637 | 1592 | 1663 |

\$seasonal

| | Jan | Feb | Mar | Apr | May | Jun |
|------|----------|-----------|----------|----------|----------|----------|
| 1994 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 1995 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 1996 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 1997 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 1998 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |

| | | | | | | |
|------|----------|-----------|----------|----------|----------|----------|
| 1999 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 2000 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 2001 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 2002 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 2003 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 2004 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |
| 2005 | 25.16919 | -82.90657 | 75.61237 | 45.65783 | 97.34343 | 16.80934 |

| | Jul | Aug | Sep | Oct | Nov | Dec |
|------|----------|-----------|-----------|-----------|-----------|----------|
| 1994 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 1995 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 1996 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 1997 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 1998 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 1999 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 2000 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 2001 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 2002 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 2003 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 2004 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |
| 2005 | 12.89646 | -13.32323 | -71.29293 | -27.92929 | -73.19066 | -4.84596 |

\$trend

| | Jan | Feb | Mar | Apr | May | Jun | Jul |
|------|----------|----------|----------|----------|----------|----------|----------|
| 1994 | NA | NA | NA | NA | NA | NA | 1363.625 |
| 1995 | 1384.042 | 1385.333 | 1386.125 | 1387.083 | 1387.833 | 1388.583 | 1389.458 |
| 1996 | 1393.917 | 1393.625 | 1394.375 | 1395.958 | 1398.333 | 1401.167 | 1403.500 |
| 1997 | 1421.208 | 1426.917 | 1431.042 | 1433.542 | 1435.375 | 1437.000 | 1439.208 |
| 1998 | 1448.208 | 1447.875 | 1448.125 | 1449.750 | 1452.583 | 1456.292 | 1460.583 |
| 1999 | 1486.750 | 1489.833 | 1494.125 | 1498.458 | 1502.500 | 1506.708 | 1511.292 |
| 2000 | 1536.875 | 1541.083 | 1543.458 | 1545.083 | 1545.917 | 1545.125 | 1543.083 |
| 2001 | 1530.958 | 1530.375 | 1530.500 | 1531.292 | 1532.833 | 1535.167 | 1538.167 |
| 2002 | 1559.667 | 1562.833 | 1565.958 | 1567.708 | 1568.875 | 1570.417 | 1572.500 |
| 2003 | 1577.375 | 1578.208 | 1578.833 | 1580.333 | 1582.125 | 1583.875 | 1585.125 |
| 2004 | 1593.083 | 1595.417 | 1597.583 | 1599.417 | 1600.917 | 1602.083 | 1603.458 |
| 2005 | 1626.917 | 1632.333 | 1638.000 | 1643.000 | 1647.792 | 1652.542 | NA |
| | Aug | Sep | Oct | Nov | Dec | | |
| 1994 | 1368.625 | 1373.250 | 1376.708 | 1379.125 | 1381.625 | | |
| 1995 | 1392.125 | 1395.000 | 1396.250 | 1396.625 | 1395.458 | | |
| 1996 | 1403.417 | 1403.875 | 1406.542 | 1410.000 | 1415.125 | | |
| 1997 | 1441.375 | 1442.792 | 1444.542 | 1446.792 | 1448.167 | | |
| 1998 | 1465.125 | 1470.417 | 1475.542 | 1480.000 | 1483.875 | | |

| | | | | | |
|------|----------|----------|----------|----------|----------|
| 1999 | 1518.083 | 1524.292 | 1527.875 | 1530.583 | 1532.958 |
| 2000 | 1538.708 | 1534.208 | 1531.917 | 1530.833 | 1531.000 |
| 2001 | 1541.625 | 1545.625 | 1550.000 | 1554.000 | 1557.250 |
| 2002 | 1574.208 | 1575.417 | 1576.167 | 1576.125 | 1576.208 |
| 2003 | 1587.250 | 1588.792 | 1588.958 | 1590.208 | 1591.542 |
| 2004 | 1603.875 | 1605.417 | 1609.583 | 1614.667 | 1620.958 |
| 2005 | NA | NA | NA | NA | NA |

\$random

| | Jan | Feb | Mar | Apr | May |
|------|--------------|--------------|--------------|--------------|--------------|
| 1994 | NA | NA | NA | NA | NA |
| 1995 | -5.21085859 | -7.42676768 | -8.73737374 | -5.74116162 | -1.17676768 |
| 1996 | -4.08585859 | 37.28156566 | -0.98737374 | -0.61616162 | -16.67676768 |
| 1997 | -13.37752525 | -16.01010101 | -6.65404040 | -5.19949495 | -3.71843434 |
| 1998 | -7.37752525 | -17.96843434 | -8.73737374 | 5.59217172 | 6.07323232 |
| 1999 | 6.08080808 | -2.92676768 | 15.26262626 | 9.88383838 | 10.15656566 |
| 2000 | 16.95580808 | 47.82323232 | 12.92929293 | 2.25883838 | -7.26010101 |
| 2001 | -7.12752525 | -16.46843434 | -7.11237374 | -5.94949495 | 1.82323232 |
| 2002 | 6.16414141 | -7.92676768 | 12.42929293 | 7.63383838 | 11.78156566 |
| 2003 | 12.45580808 | -6.30176768 | 11.55429293 | 1.00883838 | -8.46843434 |
| 2004 | 5.74747475 | 18.48989899 | -12.19570707 | -9.07449495 | -6.26010101 |
| 2005 | -9.08585859 | -27.42676768 | -6.61237374 | 1.34217172 | 14.86489899 |
| | Jun | Jul | Aug | Sep | Oct |
| 1994 | NA | 12.47853535 | 13.69823232 | 16.04292929 | 5.22095960 |
| 1995 | 15.60732323 | 11.64520202 | -3.80176768 | 7.29292929 | -4.32070707 |
| 1996 | -19.97601010 | -16.39646465 | -8.09343434 | 9.41792929 | 12.38762626 |
| 1997 | 17.19065657 | 20.89520202 | 17.94823232 | 5.50126263 | -0.61237374 |
| 1998 | 3.89898990 | -5.47979798 | -8.80176768 | -13.12373737 | -1.61237374 |
| 1999 | -7.51767677 | -26.18813131 | -17.76010101 | -7.99873737 | -8.94570707 |
| 2000 | -14.93434343 | 5.02020202 | -0.38510101 | 1.08459596 | 7.01262626 |
| 2001 | 3.02398990 | 0.93686869 | -8.30176768 | -2.33207071 | -0.07070707 |
| 2002 | -0.22601010 | -7.39646465 | 9.11489899 | -7.12373737 | -9.23737374 |
| 2003 | -4.68434343 | -1.02146465 | -2.92676768 | -6.49873737 | -0.02904040 |
| 2004 | -11.89267677 | 6.64520202 | 10.44823232 | -1.12373737 | 1.34595960 |
| 2005 | 20.64898990 | NA | NA | NA | NA |
| | Nov | Dec | | | |
| 1994 | 6.06565657 | -6.77904040 | | | |
| 1995 | -3.43434343 | -10.61237374 | | | |
| 1996 | 13.19065657 | 7.72095960 | | | |
| 1997 | -4.60101010 | -5.32070707 | | | |
| 1998 | 0.19065657 | 9.97095960 | | | |


```

1999    1.60732323    9.88762626
2000    1.35732323   -7.15404040
2001    4.19065657   -3.40404040
2002   -6.93434343    3.63762626
2003   -0.01767677    9.30429293
2004  -10.47601010   -6.11237374
2005             NA             NA

```

```
$figure
```

```

[1] 25.16919 -82.90657 75.61237 45.65783 97.34343 16.80934
[7] 12.89646 -13.32323 -71.29293 -27.92929 -73.19066 -4.84596

```

```
$type
```

```
[1] "additive"
```

```
attr(,"class")
```

```
[1] "decomposed.ts"
```

```
# Fit ARIMA model for the Milk Production dataset
```

```
arima_milk <- fit_arima(milk_data, "Monthly Milk Production")
```

Fitting ARIMA model for Monthly Milk Production

```
Warning in adf.test(ts_data, alternative = "stationary"): p-value
smaller than printed p-value
```

```
ADF Test p- value : 0.01
```

```
Series: ts_data
```

```
ARIMA(2,1,1)
```

```
Coefficients:
```

```

          ar1      ar2      ma1
      0.2066  0.3330 -0.9109
s.e.  0.0879  0.0869  0.0336

```

```
sigma^2 = 3373: log likelihood = -782.65
```

```
AIC=1573.29  AICc=1573.58  BIC=1585.14
```

```
Training set error measures:
```

```


              ME      RMSE      MAE      MPE      MAPE      MASE
Training set 10.57461 57.26184 47.71957 0.5838515 3.166587 1.536337

```

ACF1

Training set 0.002171886

Monthly Milk Production ARIMA Forecast



```
# Fit SARIMA model for the Milk Production dataset
```

```
sarima_milk <- fit_sarima(milk_data, "Monthly Milk Production")
```

Fitting SARIMA model for Monthly Milk Production

Series: ts_data

ARIMA(1,0,0)(2,1,2)[12] with drift

Coefficients:

| | ar1 | sar1 | sar2 | sma1 | sma2 | drift |
|------|--------|--------|---------|---------|--------|--------|
| | 0.8638 | 0.0607 | -0.4074 | -1.0121 | 0.4831 | 2.1882 |
| s.e. | 0.0475 | 0.1862 | 0.1173 | 0.1994 | 0.1881 | 0.2174 |

sigma² = 137.9: log likelihood = -518.84


AIC=1051.67 AICc=1052.57 BIC=1071.85

Training set error measures:

| | ME | RMSE | MAE | MPE | MAPE |
|--------------|------------|----------|----------|-------------|-----------|
| Training set | -0.1211196 | 10.98512 | 8.342375 | -0.01115387 | 0.5520753 |
| | MASE | ACF1 | | | |

Training set 0.2685838 -0.08850265

Monthly Milk Production SARIMA Forecast



Compare ARIMA and SARIMA models based on their forecast accuracy

```
compare_models(arima_milk, sarima_milk, milk_data)
```

Comparing ARIMA and SARIMA models :

ARIMA Forecast Accuracy :

21.88452 64.7399 54.04997 1.188661 3.264899 SARIMA Forecast Accuracy :

-17.81697 28.77242 19.52245 -1.093511 1.195402

Forecasting and plot comparison for Milk Production dataset

```
h_milk <- 12 # Define forecast horizon for Milk Production dataset (12 months ahead )
```

Extract the actual values for the last 12 months of the Milk Production data

```
milk_actual_values <- milk_data[(length(milk_data) - h_milk + 1): length(milk_data)]
```

Generate ARIMA forecast for the next 12 months

```
arima_milk_forecast <- forecast(arima_milk, h = h_milk)$mean
```

Generate SARIMA forecast for the next 12 months

```
sarima_milk_forecast <- forecast(sarima_milk, h = h_milk)$mean
```


Extract the time points for the last 12 months

```
time_points_milk <- time(milk_data)[(length(milk_data) - h_milk + 1): length(milk_data)]
```

Plot and compare the forecasts from ARIMA and SARIMA models against the actual values

```
plot_forecast_comparison(milk_actual_values, arima_milk_forecast, sarima_milk_forecast,  
                          time_points_milk)
```

Forecast Comparison



Program - 10

Interactive Visualization with plotly and Dynamic Reports with RMark-down

Date of Execution - 2025-10-28

Objective - This program tests students' abilities to create interactive visualizations and generate dynamic reports using plotly and RMarkdown.

```
# Load necessary libraries
library(plotly)
library(gapminder)
library(dplyr)

# Load Gapminder dataset
data("gapminder")

# Scatter plot with plotly
# Scatter plot of GDP vs Life Expectancy by Continent
scatter_plot <- gapminder %>%
  plot_ly(x = ~gdpPercap, y = ~lifeExp, color = ~continent, size = ~pop,
    hoverinfo = 'text', text = ~paste("Country:", country, "<br>GDP per Capita:", gdpPercap),
    type = 'scatter', mode = 'markers') %>%
  layout(title = 'GDP vs Life Expectancy by Continent',
    margin = list(l = 20, r = 20, b = 20, t = 30)
  )

# Display the scatter plot
scatter_plot
```


Warning: 'line.width' does not currently support multiple values.

Warning: 'line.width' does not currently support multiple values.

Warning: 'line.width' does not currently support multiple values.


Warning: 'line.width' does not currently support multiple values.

Warning: 'line.width' does not currently support multiple values.



```
# Bar chart with plotly
# Filter for year 2007 and create a bar chart of life expectancy by country
bar_chart <- gapminder %>%
  filter(year == 2007) %>%
  plot_ly(x = ~country, y = ~lifeExp, type = 'bar',
```

```
        hoverinfo = 'text', text = ~paste("Country:", country, "<br>Life Expectancy:", lifeExp)) %>%  
    layout(title = 'Life Expectancy by Country in 2007',  
           margin = list(l = 20, r = 20, b = 20, t = 30)  
    )  
  
# Display the bar chart  
bar_chart
```




```
# Line chart with plotly
# Filter data for Asia and create a line chart showing life expectancy trends over time
line_chart <- gapminder %>%
  filter(continent == 'Asia') %>%
  plot_ly(x = ~year, y = ~lifeExp, color = ~country, type = 'scatter', mode = 'lines') %>%
```



```
layout(title = 'Life Expectancy Trend in Asia',  
       margin = list(l = 20, r = 20, b = 20, t = 30)  
)  
  
# Display the line chart  
line_chart
```

Warning in RColorBrewer::brewer.pal(max(N, 3L), "Set2"): n too large, allowed maximum for palette Set2
Returning the palette you asked for with that many colors

Warning in RColorBrewer::brewer.pal(max(N, 3L), "Set2"): n too large, allowed maximum for palette Set2
Returning the palette you asked for with that many colors



```
# Combine the Plots
# Combine the scatter, bar, and line charts into one interactive layout
dashboard <- subplot(scatter_plot, bar_chart, line_chart, nrows = 1) %>%
  layout(title = 'Gapminder Data Visualization',
    margin = list(l = 20, r = 20, b = 20, t = 30))
```

```
)
```

```
Warning: 'line.width' does not currently support multiple values.
```

```
Warning: 'line.width' does not currently support multiple values.
```

```
Warning: 'line.width' does not currently support multiple values.
```

```
Warning: 'line.width' does not currently support multiple values.
```


```
Warning: 'line.width' does not currently support multiple values.
```


```
Warning in RColorBrewer::brewer.pal(max(N, 3L), "Set2"): n too large, allowed maximum for palette Set2  
Returning the palette you asked for with that many colors
```

```
Warning in RColorBrewer::brewer.pal(max(N, 3L), "Set2"): n too large, allowed maximum for palette Set2  
Returning the palette you asked for with that many colors
```

```
# Display the dashboard
```

```
dashboard
```





B.M.S. College of Engineering

Dept. of CSE (Data Science)

Basavanagudi, Bangalore-19